

1 **The statistical analysis of non-linear longitudinal data**
2 **in biomedical research using generalized additive**
3 **models**

4 *Beyond repeated measures ANOVA and Linear Mixed Models*

5 Ariel I. Mundo¹, John R. Tipton², and Timothy J. Muldoon¹

6 ¹Department of Biomedical Engineering, University of Arkansas, Fayetteville, AR, USA
7 ²Department of Mathematical Sciences, University of Arkansas, Fayetteville, AR, USA

8 **Contents**

9 **1 Abstract** **2**

10 **2 Background** **2**

11 **3 Challenges presented by longitudinal studies** **4**

12 3.1 The repeated measures ANOVA 5

13 3.2 Linear relationship 5

14 3.3 Covariance in rm-ANOVA and LMEMs 6

15 3.4 Missing observations 6

16 3.5 What does an rm-ANOVA fit looks like? A visual representation using simulated data 7

17 **4 GAMs as a special case of Generalized Linear Models** **9**

18 4.1 GAMs and Basis Functions 9

19 **5 The analysis of longitudinal biomedical data using GAMs** **12**

20 5.1 Simulated data 12

21 5.2 An interaction GAM for longitudinal data 12

22 5.3 Determination of significance in GAMs for longitudinal data 14

23 **6 Discussion** **15**

24 **7 Conclusion** **16**

25 **8 Acknowledgements** **16**

26 **9 References** **17**

27	A Code for Manuscript data	20
28	A.1 Compound symmetry and independent errors in linear and quadratic responses	21
29	A.2 Basis functions and GAMs	28
30	B Longitudinal biomedical data simulation and GAMs	32
31	B.1 A basic Workflow for GAMs	34
32	C GAM and Linear model plots and Missing data	44
33	C.1 GAM and Linear model plots	45
34	C.2 Working with Missing data in GAMs	46
35	C.3 Pairwise comparisons in GAMs: full and missing data cases	50

36 1 Abstract

37 In biomedical research, the outcome of longitudinal studies has been traditionally analyzed using the *re-*
38 *peated measures analysis of variance* (rm-ANOVA) or more recently, *linear mixed models* (LMEMs). Al-
39 though LMEMs are less restrictive than rm-ANOVA in terms of correlation and missing observations, both
40 methodologies share an assumption of linearity in the measured response, which results in biased estimates
41 and unreliable inference when they are used to analyze data where the trends are non-linear. In contrast,
42 generalized additive models (GAMs) relax the linearity assumption, and allow the data to determine the fit
43 of the model while permitting missing observations and different correlation structures. Therefore, GAMs
44 present an excellent choice to analyze non-linear longitudinal data in the context of biomedical research.
45 This paper summarizes the limitations of rm-ANOVA and LMEMs, presents the basic theory of GAMs, and
46 uses simulated data that follows trends reported in the biomedical literature to demonstrate their implemen-
47 tation in R via the package *mgcv*. To make this work reproducible, the code and data used in this paper are
48 available at:_____.

49 2 Background

50 Longitudinal studies are designed to repeatedly measure a variable of interest in a group (or groups) of
51 subjects, with the intention of observing the evolution of effect across time rather than analyzing a single
52 time point (e.g., a cross-sectional study). Biomedical research frequently uses longitudinal studies to analyze
53 the evolution of a “treatment” effect across multiple time points; and in such studies the subjects of analysis
54 range from animals (mice, rats, rabbits), to human patients, cells, or blood samples, among many others.
55 Tumor response [1–4], antibody expression [5,6], and cell metabolism [7,8] are examples of the different
56 situations where researchers have used longitudinal designs to study some physiological response. Because
57 the frequency of the measurements in a longitudinal study is dependent on the biological phenomena of
58 interest and the experimental design of the study, the frequency of such measurements can range from minute
59 intervals to study a short-term response such as anesthesia effects in animals[9], to weekly measurements
60 to analyze a mid-term response like the evolution of dermatitis symptoms in breast cancer patients [10], to
61 monthly measurements to study a long-term response such as mouth opening following radiotherapy (RT)
62 in neck cancer patients [11].

63 Traditionally, a “frequentist” or “classical” statistical paradigm is used in biomedical research to derive
64 inferences from a longitudinal study. The frequentist paradigm regards probability as the limit of the
65 expected outcome when an experiment is repeated a large number of times [12], and such view is applied
66 to the analysis of longitudinal data by assuming a null hypothesis under a statistical model that is often an
67 *analysis of variance over repeated measures* (repeated measures ANOVA or rm-ANOVA). The rm-ANOVA

model makes three key assumptions regarding longitudinal data: 1) linearity of the response across time, 2) constant correlation across same-subject measurements, and 3) observations from each subject are obtained at all time points through the study (a condition also known as *complete observations*) [13,14].

The expected linear behavior of the response through time is a key requisite in rm-ANOVA [15]. This “linearity assumption” in rm-ANOVA implies that the model is misspecified when the data does not follow a linear trend, which results in unreliable inference. In biomedical research, non-linear trends are the norm rather than the exception in longitudinal studies. A particular example of this non-linear behavior in longitudinal data arises in measurements of tumor response to chemo and/or radiotherapy in preclinical and clinical settings [1,8,16]. These studies have shown that the collected signal does not follow a linear trend over time, and presents extreme variability at different time points, making the fit of rm-ANOVA model inconsistent with the observed variation. Therefore, when rm-ANOVA is used to draw inference of such data the estimates are inevitably biased, because the model is only able to accommodate linear trends that fail to adequately represent the biological phenomenon of interest.

A *post hoc* analysis is often used in conjunction with rm-ANOVA to perform repeated comparisons to estimate a *p-value*, which in turn is used as a measure of significance. Although it is possible that a *post hoc* analysis of rm-ANOVA is able to find “significant” *p-values* ($p < 0.05$) from non-linear data, the validity of such metric is dependent on how adequate the model fits the data. In other words, *p-values* are valid only if the model and the data have good agreement; if that is not the case, a “Type III” error (known as “model misspecification”) occurs [17]. For example, model misspecification will occur when a model that is only able to explain linear responses (such as rm-ANOVA) is fitted to data that follows a quadratic trend, thereby causing the resulting *p-values* and parameter estimates to be invalid [18].

Additionally, the *p-value* itself is highly variable, and multiple comparisons can inflate the false positivity rate (Type I error or α) [19,20], consequently biasing the conclusions of the study. Corrections exist to address the Type I error issue of multiple comparisons (such as Bonferroni [21]), but they in turn reduce statistical power ($1 - \beta$) [22], and lead to increased Type II error (failing to reject the null hypothesis when the null hypothesis is false) [23,24]. Therefore, the tradeoff of *post hoc* comparisons in rm-ANOVA between Type I, II and III errors might be difficult to resolve in a biomedical longitudinal study where a delicate balance exists between statistical power and sample size.

On the other hand, the assumption of constant correlation in rm-ANOVA (often known as the *compound symmetry assumption*) is typically unreasonable because correlation between the measured responses often diminishes as the time interval between the observation increases [25]. Corrections can be made in rm-ANOVA in the absence of compound symmetry [26,27], but the effectiveness of the correction is limited by the size of the sample, the number of measurements [28], and group sizes [29]. In the case of biomedical research, where living subjects are frequently used, sample sizes are often not “large” due to ethical and budgetary reasons [30] which might cause the corrections for lack of compound symmetry to be ineffective.

Due to a variety of causes, the number of observations during a study can vary between all subjects. For example, in a clinical trial patients may voluntarily withdraw, whereas attrition due to injury or weight loss in preclinical animal studies is possible. It is even plausible that unexpected complications with equipment or supplies arise that prevent the researcher from collecting measurements at certain time points. In each of these missing data scenarios, the *complete observations* assumption of classical rm-ANOVA is violated. When incomplete observations occur, a rm-ANOVA model is fit by excluding all subjects with missing observations from the analysis [13]. This elimination of partially missing data from the analysis can result in increased costs if the desired statistical power is not met with the remaining observations, because it would be necessary to enroll more subjects. At the same time, if the excluded observations contain insightful information that is not used, their elimination from the analysis may limit the demonstration of significant differences between groups.

During the last decade, the biomedical community has started to recognize the limitations of rm-ANOVA in the analysis of longitudinal data. The recognition on the shortcomings of rm-ANOVA is exemplified by the use of linear mixed effects models (LMEMs) by certain groups to analyze longitudinal tumor response data [8,16]. Briefly, LMEMs incorporate *fixed effects*, which correspond to the levels of experimental factors in the study (e.g., the different drug regimens in a clinical trial), and *random effects*, which account for random

variation within the population (e.g., the individual-level differences not due to treatment such as weight or age). When compared to the traditional rm-ANOVA, LMEMs are more flexible as they can accommodate missing observations for multiple subjects and allow different modeling strategies for the variability within each measure in every subject [15]. However, LMEMs impose restrictions in the distribution of the errors of the random effects, which need to be normally distributed and independent [13,31]. And even more importantly, LMEMs also assume a linear relationship between the response and time [15], making them unsuitable to analyze non-linear data.

As the rm-ANOVA and the more flexible LMEM approaches make overly restrictive assumptions regarding the linearity of the response, there is a need for biomedical researchers to explore the use of additional statistical tools that allow the data (and not an assumption in trend) to determine the trend of the fitted model, to enable appropriate inference.

In this regard, generalized additive models (GAMs) present an alternative approach to analyze longitudinal data. Although not frequently used by the biomedical community, these semi-parametric models are customarily used in other fields to analyze longitudinal data. Examples of the use of GAMs include the analysis of temporal variations in geochemical and palaeoecological data [32–34], health-environment interactions [35] and the dynamics of government in political science [36]. There are several advantages of GAMs over LMEMs and rm-ANOVA models: 1) GAMs can fit a more flexible class of smooth responses that enable the data to dictate the trend in the fit of the model, 2) they can model non-constant correlation between repeated measurements [37] and 3) can easily accommodate missing observations. Therefore, GAMs can provide a more flexible statistical approach to analyze non-linear biomedical longitudinal data than LMEMs and rm-ANOVA.

The current advances in programming languages designed for statistical analysis (specifically R), have eased the computational implementation of traditional models such as rm-ANOVA and more complex approaches such as LMEMs and GAMs. In particular, R[38] has an extensive collection of documentation and functions to fit GAMs in the package *mgcv* [37,39] that not only speed up the initial stages of the analysis but also enable the use of advanced modeling structures (e.g. hierarchical models, confidence interval comparisons) without requiring advanced programming skills from the user. At the same time, R has many tools that simplify data simulation, an emerging strategy used to test statistical models [28]. Data simulation methods allow the researcher to create and explore different alternatives for analysis without collecting information in the field, reducing the time window between experiment design and its implementation, and simulation can be also used for power calculations and study design questions.

This work provides biomedical researchers with a clear understanding of the theory and the practice of using GAMs to analyze longitudinal data using by focusing on four areas. First, the limitations of LMEMs and rm-ANOVA regarding linearity of response, constant correlation structures and missing observations is explained in detail. Second, the key theoretical elements of GAMs are presented using clear and simple mathematical notation while explaining the context and interpretation of the equations. Third, using simulated data that reproduces patterns in previously reported studies [16] we illustrate the type of non-linear longitudinal data that often occurs in biomedical research. The simulated data experiments highlight the differences in inference between rm-ANOVA, LMEMs and GAMs on data similar to what is commonly observed in biomedical studies. Finally, reproducibility is emphasized by providing the code to generate the simulated data and the implementation of different models in R, in conjunction with a step-by-step guide demonstrating how to fit models of increasing complexity.

In summary, this work will allow biomedical researchers to identify when the use of GAMs instead of rm-ANOVA or LMEMs is appropriate to analyze longitudinal data, and provide guidance on the implementation of these models by improving the standards for reproducibility in biomedical research.

3 Challenges presented by longitudinal studies

3.1 The repeated measures ANOVA

The *repeated measures analysis of variance* (rm-ANOVA) is the standard statistical analysis for longitudinal data in biomedical research. This statistical methodology requires certain assumptions for the model to be valid. From a practical view, the assumptions can be divided in three areas: 1) linear relationship between covariates and response, 2) a constant correlation between measurements, and, 3) complete observations for all subjects. Each one of these assumptions is discussed below.

3.2 Linear relationship

3.2.1 The repeated measures ANOVA case

In a longitudinal biomedical study, two or more groups of subjects (e.g., human subject, mice, samples) are subject to different treatments (e.g., a “treatment” group receives a novel drug or intervention vs. a “control” group that receives a placebo), and measurements from each subject within each group are collected at specific time points. The collected response is modeled with *fixed* components. The *fixed* component can be understood as a constant value in the response which the researcher is interested in measuring, i.e., the average effect of the novel drug/intervention in the “treatment” group.

Mathematically speaking, a rm-ANOVA model with an interaction can be written as:

$$y_{ijt} = \beta_0 + \beta_1 \times time_t + \beta_2 \times treatment_j + \beta_3 \times time_t \times treatment_j + \varepsilon_{ijt} \quad (1)$$

In this model y_{ijt} is the response for subject i , in treatment group j at time t , which can be decomposed in a mean value β_0 , *fixed effects* of time ($time_t$), treatment ($treatment_j$) and their interaction $time_t * treatment_j$ which have linear slopes given by β_1, β_2 and β_3 , respectively. Independent errors ε_{tij} represent random variation not explained by the *fixed* effects, and are assumed to be $\sim N(0, \sigma^2)$ (independently and identically normally distributed with mean zero and variance σ^2). In a biomedical research context, suppose two treatments groups are used in a study (e.g., “placebo” vs. “novel drug” or “saline” vs. “chemotherapy”). Then, the group terms in Equation (1) can be written as below with $treatment_j = 0$ representing the first treatment group (Group A) and $treatment_j = 1$ representing the second treatment group (Group B). The linear models then can be expressed as

$$y_{ijt} = \begin{cases} \beta_0 + \beta_1 \times time_t + \varepsilon_{ijt} & \text{if Group A} \\ \beta_0 + \beta_2 + \beta_1 \times time_t + \beta_3 \times time_t + \varepsilon_{ijt} & \text{if Group B} \end{cases} \quad (2)$$

To further simplify the expression, substitute $\widetilde{\beta}_0 = \beta_0 + \beta_2$ and $\widetilde{\beta}_1 = \beta_1 + \beta_3$ in the equation for Group B. This substitution allows for a different intercept and slope for Groups A and B. The model is then written as

$$y_{ijt} = \begin{cases} \beta_0 + \beta_1 \times time_t + \varepsilon_{ijt} & \text{if Group A} \\ \widetilde{\beta}_0 + \widetilde{\beta}_1 \times time_t + \varepsilon_{ijt} & \text{if Group B} \end{cases} \quad (3)$$

Presenting the model in this manner makes clear that when treating different groups, an rm-ANOVA model is able to accommodate non-parallel lines in each case (different intercepts and slopes per group). In other words, the rm-ANOVA model “expects” a linear relationship between the covariates and the response, this means that either presented as Equation (1), Equation (2) or Equation (3), an rm-ANOVA model is only able to accommodate linear patterns in the data. If the data show non-linear behavior, the rm-ANOVA model will approximate this behavior with non-parallel lines.

3.2.2 The Linear Mixed Model Case

A linear mixed model (LMEM) is a class of statistical model that incorporates *fixed effects* to model the relationship between the covariates and the response, and *random effects* to model subject variability that is not the primary focus of the study but that might be important to distinguish [15,40]. A LMEM with interaction between time and treatment for a longitudinal study can be written as:

$$y_{ijt} = \beta_0 + \beta_1 \times time_t + \beta_2 \times treatment_j + \beta_3 \times time_t \times treatment_j + \mu_{ij} + \varepsilon_{ijt} \quad (4)$$

When Equation (1) and Equation (4) are compared, it is easily noticeable that LMEM and rm-ANOVA have the same construction regarding the *fixed effects* of time and treatment, but that the LMEM incorporates an additional source of variation (the term μ_{ij}). This term μ_{ij} is the one that corresponds to the *random effect*, accounting for variability in each subject within each group. The *random* component can also be understood as used to model some “noise” in the response, but that is intended to be analyzed and disentangled from the “global noise” term ε_{ijt} from Equation (1).

For example, if the blood concentration of the drug is measured in certain subjects in the early hours of the morning while other subjects are measured in the afternoon, it is possible that the difference in the collection time introduces some “noise” in the data. As the name suggests, this “random” variability needs to be modeled as a variable rather than as a constant value. The *random effect* μ_{ij} in Equation (4) is assumed to be $\mu_{ij} \sim N(0, \sigma_\mu^2)$. In essence, the *random effect* in a LMEM enables to fit models with different slopes at the subject-level [15]. However, the expected linear relationship of the covariates and the response in Equation (1) and in Equation (4) is essentially the same, representing a major limitation of LMEMs to fit a non-linear response.

3.3 Covariance in rm-ANOVA and LMEMs

In a longitudinal study there is an expected *covariance* between repeated measurements on the same subject, and because repeated measures occur in the subjects within each group, there is a *covariance* between measurements at each time point within each group. The *covariance matrix* (also known as the variance-covariance matrix) is a matrix that captures the variation between and within subjects in a longitudinal study [41] (For an in-depth analysis of the covariance matrix see [40,42]).

In the case of an rm-ANOVA analysis, it is typically assumed that the covariance matrix has a specific construction known as *compound symmetry* (also known as “sphericity” or “circularity”). Under this assumption, the between-subject variance and within-subject correlation are constant across time [26,42,43]. However, it has been shown that this condition is frequently not justified because the correlation between measurements tends to change over time [44]; and it is higher between consecutive measurements [13,25]. Although corrections can be made (such as Huynh-Feldt or Greenhouse-Geisser) [26,27] the effectiveness of each correction is limited because it depends on the size of the sample, the number of repeated measurements [28], and they are not robust if the group sizes are unbalanced [29]. Because biomedical longitudinal studies are often limited in sample size and can have an imbalanced design, the corrections required to use an rm-ANOVA model may not be able to provide a reasonable adjustment that makes the model valid.

In the case of LMEMs, one key advantage over rm-ANOVA is that they allow different structures for the variance-covariance matrix including exponential, autoregressive of order 1, rational quadratic and others [15]. Nevertheless, the analysis required to determine an appropriate variance-covariance structure for the data can be a challenging process by itself. Overall, the spherical assumption for rm-ANOVA may not capture the natural variations of the correlation in the data, and can bias the inferences from the analysis.

3.4 Missing observations

Missing observations are an issue that arises frequently in longitudinal studies. In biomedical research, this situation can be caused by reasons beyond the control of the investigator [45]. Dropout from patients

and attrition or injury in animals are among the reasons for missing observations. Statistically, missing information can be classified as *missing at random* (MAR), *missing completely at random* (MCAR), and *missing not at random* (MNAR) [42]. In a MAR scenario, the pattern of the missing information is related to some variable in the data, but it is not related to the variable of interest [46]. If the data are MCAR, this means that the missingness is completely unrelated to the collected information [47], and in the case of MNAR the missing values are dependent on their value.

An rm-ANOVA model assumes complete observations for all subjects, and therefore subjects with one or more missing observations are excluded from the analysis. This is inconvenient because the remaining subjects might not accurately represent the population, and statistical power is affected by this reduction in sample size [48]. In the case of LMEMs, inferences from the model are valid when missing observations in the data exist that are MAR or MCAR [40]. For example, if attrition occurs in all mice that had lower weights at the beginning of a chemotherapy response study, the missing data can be considered MAR because the missingness is unrelated to other variables of interest.

3.5 What does an rm-ANOVA fit look like? A visual representation using simulated data

To visually demonstrate the limitations of rm-ANOVA and LMEMs for non-linear longitudinal data, this section presents a simulation experiment of a normally distributed response of two groups of 10 subjects each. An rm-ANOVA model (Equation (1)), and a LMEM (Equation (4)) are fitted to each group, using R[38] and the package *nlme*[49].

Briefly, two cases for the mean responses for each group are considered: in the first case, the mean response in each group is a linear function with different intercepts and slopes; a negative slope is used for Group 1 and a positive slope is used for Group 2 (Figure 1, A). In the second case, a second-degree polynomial (quadratic) function is used for the mean response per group: the quadratic function is concave down for Group 1 and it is concave up for Group 2 (Figure 1, C). In both the linear and quadratic simulated data, the groups start with the same mean value at the first time point. This is intentional in order to simulate the expected temporal evolution of some physiological quantity, which is typical in biomedical experiments.

Specifically, the rationale for the chosen linear and quadratic functions is the likelihood that a measured response in two treatment groups is similar in the initial phase of the study, but as therapy progresses a divergence in the trend of the response indicates a treatment effect. In other words, Group 1 can be thought as a “Control” group and Group 2 as a “Treatment” group. From the mean response per group (linear or quadratic), the variability or “error” of individual responses within each group is simulated using a covariance matrix with compound symmetry (constant variance across time). Thus, the response per subject in both the linear and quadratic simulation corresponds to the mean response per group plus the error (Figure 1 B,D).

A more comprehensive exploration of the fit of rm-ANOVA and LMEMs for linear and non-linear longitudinal appears in Figure 5 and Figure 6 in the Appendix, where simulation with compound symmetry and independent errors (errors generated from a normal distribution that are not constant over time) and the plot of simulated errors, and fitted parameters is presented. We are aware that the simulated data used in this section present an extreme case that might not occur frequently in biomedical research, but they are used as a representation of the consequences of modeling non-linear data with a linear model such as rm-ANOVA or LMEMs. Of notice, Section 5 uses simulated data that does follow reported trends in the biomedical literature.

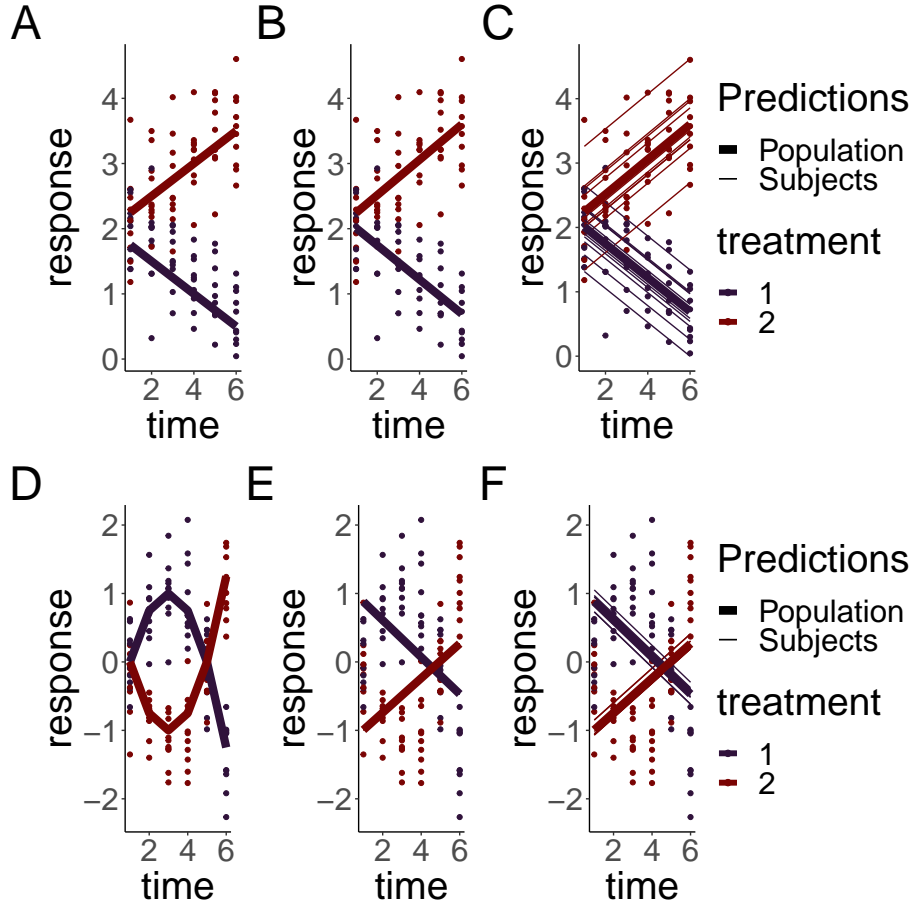


Figure 1: Simulated linear responses from two groups with correlated (top row) or independent (bottom row) errors using a LMEM and a rm-ANOVA model. A, D: Simulated data with known mean response (linear or quadratic, thin lines) and individual responses (points) showing the dispersion of the data. B,E: Estimates from the rm-ANOVA model for the mean group response (linear or quadratic). Points represent the original raw data. The rm-ANOVA model not only fails to pick the trend of the quadratic data but also assigns a global estimate that does not take between-subject variation. C, F: Estimates from the LMEM model in the linear and quadratic case. The LMEM incorporates a random effect for each subject, but this model and the rm-ANOVA model are unable to follow the trend of the data in each group and grossly bias the initial estimates for each group.

The simulation shows that the fit produced by the LMEM and the rm-ANOVA model is good for linear data, as the predictions for the mean response are reasonably close to the “truth” of the simulated data (Figure 1, B, E). When the linearity and compound symmetry assumptions are met, the model approximates well the individual trends and the mean trends by group.

However, consider the case when the data follows a non-linear trend, such as the simulated data in Figure 1, C. Here, the mean response per group was simulated using a quadratic function but errors, individual responses and the rm-ANOVA model were produced in the same manner as in (Figure 1 A and B). The mean response in the simulated data with quadratic behavior is changing in each group through the timeline, and the mean value is the same as the initial value by the fifth time point for each group. Fitting an rm-ANOVA model (1) or a LMEM (4) to this data produces the fit that appears in panels E and F in Figure 1.

A comparison of the fitted mean response of the LMEM and the rm-ANOVA model to the simulated data in Figure ((1, E, F) indicates that the models are not capturing the changes within each group. Specifically, note that the fitted mean response of both models (panel E, F) show that the change (increase for Treatment 1 or decrease for Treatment 2) in the response through time points 2 and 4 is not being captured. The LMEM

is only able to account for between-subject variation by providing different intercepts to each subject, but both models are not able to capture the fact that the initial values are the same in each group, and instead fit non-parallel lines that have initial values that are markedly different from the “true” initial values in each case (compare panel D with panels E and F). If such a change has important physiological implications, both rm-ANOVA and LMEMs omit it from the fitted mean response. Thus, even though the model correctly detects a divergence between treatment groups, the exact nature of this difference is not correctly identified, limiting valuable inferences from the data.

This section has used simulation to better convey the limitations of linearity and correlation in the response in non-linear data. Although the model fitted to the simulated data was an rm-ANOVA model, the main issue of an expected linear trend in the response is the same in the case of a LMEM. In the following section, we present generalized additive models (GAMs) as a data-driven alternative method to analyze longitudinal non-linear data.

4 GAMs as a special case of Generalized Linear Models

4.1 GAMs and Basis Functions

Generalized linear models (GLMs) are a family of models that fit a linear response function to data that do not have normally distributed errors[50]. In contrast, GAMs are a family of regression-based methods for estimating smoothly varying trends and are a broader class of models that contain the GLM family as a special case[34,37,51]. A GAM model can be written as:

$$y_{ijt} = \beta_0 + f(x_t | \beta_j) + \varepsilon_{ijt} \quad (5)$$

Where y_{ijt} is the response at time t of subject i in group j , β_0 is the expected value at time 0, the change of y_{ijt} over time is represented by the *smooth function* $f(x_t | \beta_j)$ with inputs as the covariates x_t and parameters β_j , and ε_{ijt} represents the residual error.

In contrast to the linear functions used to model the relationship between the covariates and the response in rm-ANOVA or LMEM, GAMs use more flexible *smooth functions*. This approach is advantageous as it does not restrict the model to a linear relationship, although a GAM will estimate a linear relationship if the data is consistent with a linear response. One possible set of functions for $f(x_t | \beta_j)$ that allow for non-linear responses are polynomials, but a major limitation is that polynomials create a “global” fit as they assume that the same relationship exists everywhere, which can cause problems with inference [36]. In particular, polynomial fits are known to show boundary effects because as t goes to $\pm\infty$, $f(x_t | \beta_j)$ goes to $\pm\infty$ which is almost always unrealistic, and causes bias at the endpoints of the time period.

The smooth functional relationship between the covariates and the response in GAMs is specified using a semi-parametric relationship that can be fit within the GLM framework, by using *basis function* expansions of the covariates and by estimating random coefficients associated with these basis functions. A *basis* is a set of functions that spans the mathematical space where the smooths that approximate $f(x_t | \beta_j)$ exist [34]. For the linear model in Equation (1), the basis coefficients are β_1 , β_2 and β_3 and the basis vectors are $time_t$, $treatment_j$ and $time_t \times treatment_j$. The basis function then, is the combination of basis coefficients and basis vectors that map the possible relationship between the covariates and the response [52], which in the case of Equation (1) is restricted to a linear family of functions. In the case of Equation (5), the basis function is $f(x_t | \beta_j)$, which means that the model allows for non-linear relationships among the covariates.

Commonly used *basis functions* are splines (cubic, thin plate regression among others). A cubic spline is a smooth curve constructed from cubic polynomials joined together in a manner that enforces smoothness, and thin plate regression splines are an optimized version that work well with noisy data [34,37]. Splines have a long history in solving semi-parametric statistical problems and are often a default choice to fit GAMs as they are a simple, flexible and powerful option to obtain smoothness [53]. Therefore, this data-driven

flexibility in GAMs overcomes the limitation that occurs in LMEMs and rm-ANOVA when the data is non linear.

To further clarify the concept of basis functions and smooth functions, consider the simulated response for Group 1 in Figure (1, C). The simplest GAM model that can be used to estimate such response is that of a single smooth term for the time effect; i.e., a model that fits a smooth to the trend of the group through time. The timeline can be divided in equally spaced *knots*, each knot being a region where a different basis function will be used. Because there are six timepoints for this group, five knots can be used. The model with five knots to construct the smooth term means that it will have four basis functions (plus one that corresponds to the intercept). The choice of basis functions is already optimized in the package *mgcv* depending on the number of knots. In Panel A of Figure 2, the four basis functions (and the intercept) are shown. Each of the basis functions is composed of six different points (because there are six points on the timeline). To control the “wigliness” of the fit, each of the basis functions of Panel A is weighted by multiplying it by a coefficient according to the matrix of Panel B. The parameter estimates are penalized where the penalty reduces the “wigliness” of the smooth fit to prevent overfitting: A weak penalty estimate will result in wiggly functions whereas a strong penalty estimate provides evidence that a linear response is appropriate.

In other words, the six points of each basis are multiplied by the corresponding coefficient in panel B, thereby increasing or decreasing the original basis functions of Panel A. In Figure 2, Panel C shows the resulting penalized basis functions. Note that the penalization for basis 1 has resulted in a decrease of its overall value (because the coefficient for that basis function is negative and less than 1); on the other hand, basis 3 has roughly doubled its value. Finally, the penalized basis functions are added at each timepoint to produce the smooth term. The resulting smooth term for the effect of *time* is shown in Panel D (orange line) along the simulated values per group, which appear as points.

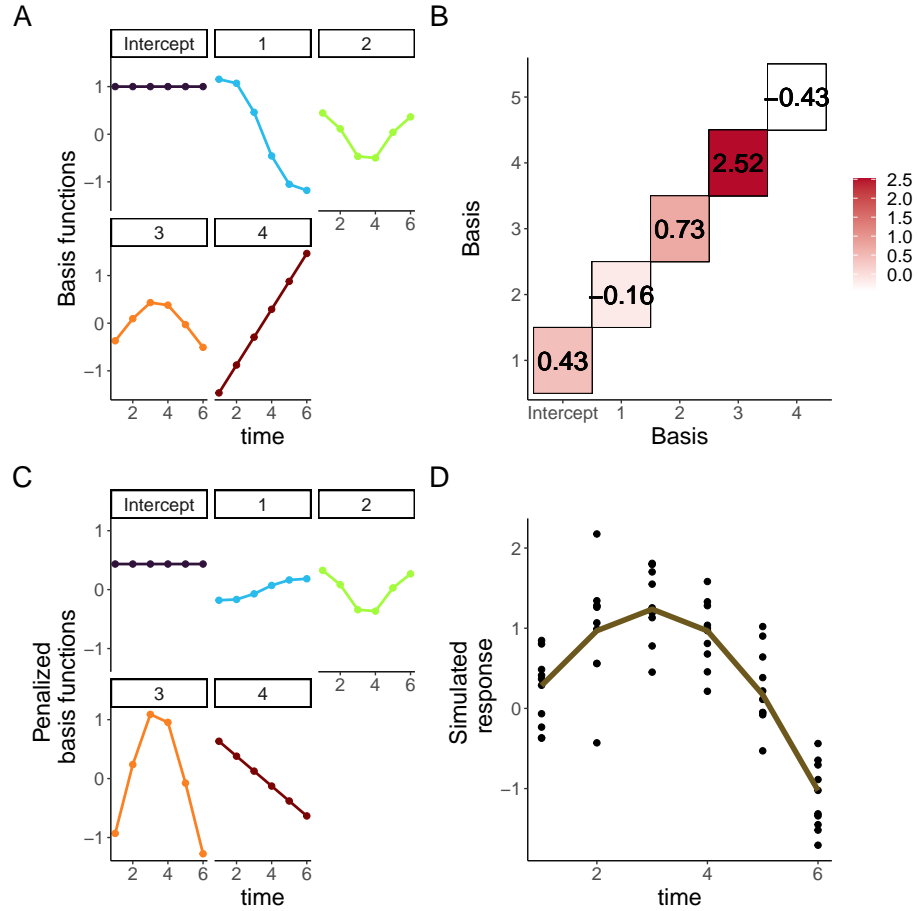


Figure 2: Basis functions for a single smoother for time with five knots. A: Basis functions for a single smoother for time for the simulated data of Group 1 from Figure 2, the intercept basis is not shown. B: Matrix for basis function penalization. Each basis function is multiplied by a coefficient which can be positive or negative. The coefficient determines the overall effect of each basis in the final smoother. C: Weighted basis functions. Each of the four basis functions of panel A has been weighted by the corresponding coefficient shown in Panel B, note the corresponding increase (or decrease) of each basis. D: Smoother for time and original data points. The smoother (line) is the result of the sum of each penalized basis function at each time point, with simulated values for the group shown as points.

5 The analysis of longitudinal biomedical data using GAMs

The previous sections provided the basic framework to understand the GAM framework and how these models are more advantageous to analyze non-linear longitudinal data when compared to rm-ANOVA or LMEMs. This section will use simulation to present the practical implementation of GAMs for longitudinal biomedical data using R and the package `mgcv`. The code for the simulated data and figures, and a brief guide for model selection and diagnostics appear in the Appendix.

5.1 Simulated data

The simulated data is based on the reported longitudinal changes in oxygen saturation (StO_2) in subcutaneous tumors that appear in Figure 3, C in [16]. In the paper, diffuse reflectance spectroscopy was used to quantify StO_2 changes in both groups at the same time points (days 0, 2, 5, 7 and 10). In the “Treatment” group (chemotherapy) an increase in StO_2 is observed through time, while a decrease is seen in the “Control” (saline) group. Following the reported trend, we simulated 10 normally distributed observations at each time point with a standard deviation (SD) of 10% (matching the SD in the original paper). The simulated and real data appear in Figure 3, A and the inset, respectively.

5.2 An interaction GAM for longitudinal data

An interaction effect is typically the main interest in longitudinal biomedical data, as it takes into account treatment, time, and their combination. In a practical sense, when a GAM is implemented for longitudinal data, a smooth can be added to the model for the *time* effect to account for the repeated measures over time. Although specific methods of how GAMs model correlation structures is a topic beyond the scope of this paper, it suffices to say that GAMs are flexible and can handle correlation structures beyond compound symmetry. A detailed description on basis functions and correlations can be found in [52].

For the data in Figure 3, A the main effect of interest is how StO_2 changes over time for each treatment. To estimate this, the model needs to incorporate independent smooths for *Group* and *Day*, respectively. The main thing to consider is that model syntax accounts for the fact that one of the variables is numeric (*Day*) and the other is a factor (*Group*). Because the smooths are centered at 0, the factor variable needs to be specified as a parametric term in order to identify any differences between the groups. Using R and the package `mgcv` the model syntax is:

```
m1<-gam(StO2_sim~Group+s(Day,by=Group,k=5), method='REML',data=dat_sim)
```

This syntax specifies that `m1` will store the model, and that the change in the simulated oxygen saturation (`StO2_sim`) is modeled using independent smooths for *Group* and *Day* (the parenthesis preceded by *s*) using 5 knots. The smooth is constructed by default using thin plate regression splines. Other splines can be used if desired, including gaussian process smooths [34]. The parametric term *Group* is added to quantify differences in the effect of treatment between groups, and the `method` chosen to select the smoothing parameters is the restricted maximum likelihood (REML) [37]. When the smooths are plotted over the raw data, it is clear that the model has been able to capture the trend of the change of StO_2 for each group across time (Figure 3,B). Model diagnostics can be obtained using the `gam.check` function, and the function `appraise` from the package *gratia* [54]. A guide for model selection and diagnostics is in the Appendix, and an in-depth analysis can be found in [37] and [55].

One question that might arise at this point is “what is the fit that an rm-ANOVA model produces for the simulated data?” The rm-ANOVA model, which corresponds to Equation (1) is presented in Figure 3,C. This is a typical case of model misspecification: The slopes of each group are different, which would lead to a *p-value* indicating significance for the treatment and time effects, but the model is not capturing the changes that occur at days 2 and between days 5 and 7, whereas the GAM model is able to do so (Figure 3,B) .

406 Because GAMs do not require equally-spaced or complete observations for all subjects, they are advantageous
 407 to analyze longitudinal data where missingness exists. The rationale behind this is that GAMs are able to
 408 pick the trend in the data even when some observations are missing. However, this usually causes the
 409 resulting smooths to have wider confidence intervals and less ability to pick certain trends. Consider the
 410 simulated StO_2 values from Figure (3, B). If 40% of the total observations are randomly deleted and the
 411 same interaction GAM fitted for the complete dataset is used, the resulting smooths are still able to show a
 412 different trend for each group, but it can be seen that the smooths overlap during the first 3 days because
 413 with less data points, the trend is less pronounced than in the full dataset (3, D). Although the confidence
 414 intervals have increased for both smooths, the model still shows different trends with as little as 4 observations
 415 per group at certain time points.

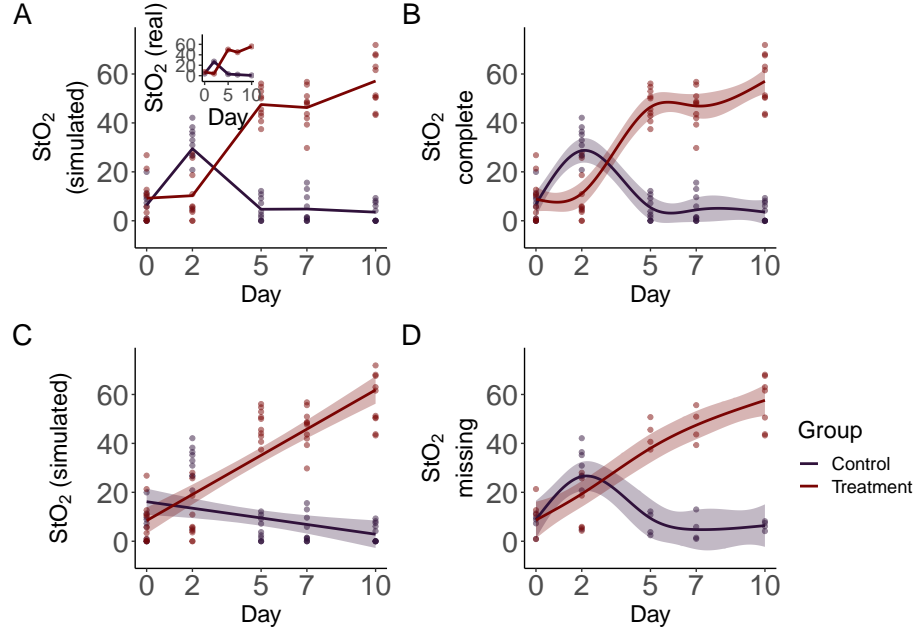


Figure 3: Simulated data and smooths for oxygen saturation in tumors. A: Simulated data that follows previously reported trends (inset) in tumors under chemotherapy (Treatment) or saline (Control) treatment. Simulated data is from a normal distribution with standard deviation of 10% with 10 observations per time point. Lines indicate mean oxygen saturation B: Smooths from the GAM model for the full simulated data with interaction of Group and Treatment. Lines represent trends for each group, shaded regions are 95% confidence intervals. C: rm-ANOVA model for the simulated data, the model does not capture the changes in each group over time. D: Smooths for the GAM model for the simulated data with 40% of its observations missing. Lines represent trends for each group, shaded regions are 95% confidence intervals.

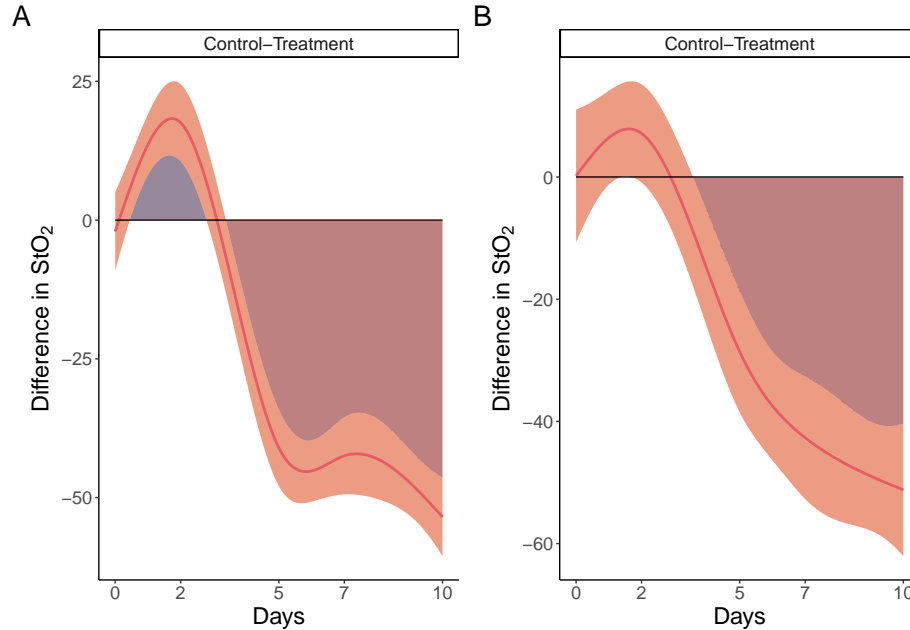


Figure 4: Pairwise comparisons for smooth terms. A: Pairwise comparisons for the full dataset. B: Pairwise comparisons for the dataset with missing observations. Significant differences exist where the interval does not cover 0. In both cases the effect of treatment is significant after day 5.

5.3 Determination of significance in GAMs for longitudinal data

At the core of a biomedical longitudinal study lies the question of a significant difference between the effect of two or more treatments in different groups. Whereas in *rm*-ANOVA a *post-hoc* analysis is required to answer such question by calculating some *p-values* after multiple comparisons, GAMs can use a different approach to estimate significance. In essence, the idea behind the estimation of significance in GAMs across different treatment groups is that if the *difference* between the confidence intervals of the fitted smooths for such groups is non-zero, then a significant difference exists at that time point(s). The absence of a *p-value* in this case might seem odd, but the confidence interval comparison can be conceptualized in the following manner: Different trends in each group are an indication of an effect by the treatment. This is what happens for the simulated data in Figure 3, A where the chemotherapy causes StO_2 to increase over time.

With this expectation of different trends in each group, computing the difference between the trends will identify if the observed change is significant. The difference between groups with similar trends is likely to yield zero, which would indicate that the treatment is not causing a change in the response in one of the groups (assuming the other group is a Control or Reference group).

Consider the calculation of pairwise differences for the smooths in Figure 3, B and D. Figure 4, shows the comparison between each treatment group for the full and missing datasets. Here, the “Control” group is used as the reference to which “Treatment” group is being compared. Of notice, the pairwise comparison has been set on the response scale (see Appendix for code details), because otherwise the comparison appears shifted and is not intuitively easy to relate to the original data.

With this correction in mind, the shaded region under or above the confidence interval (that does not cover 0) indicates the interval where each group has a higher effect than the other. Notice that the shaded region between days 0 and 3 for the full dataset indicates that through that time, the “Control” group has higher StO_2 , but as therapy progresses the effect is reversed and by day 4 it is the “Treatment” group the one that has greater StO_2 . This would suggest that the effect of chemotherapy in the “Treatment” group becomes significant after day 4 for the model used. Moreover, notice that although there is no actual measurement at day 4, the model is capable of providing an estimate of when the shift in StO_2 occurs.

On the data with missing observations (Figure 3, D), the confidence intervals of the smooths overlap between days 0 and 3. Consequently, the smooth pairwise comparison (Figure 4, B) shows that there is not a significant difference between the groups during that period, but is still able to pick the change on day 4 as the full dataset smooth pairwise comparison.

In a sense, the pairwise smooth comparison is more informative than a *post-hoc p-value*. For biomedical studies, it is able to provide an estimate of *when* a biological process becomes significant. This is advantageous because it can help researchers gain insight on metabolic changes and other biological processes that can be worth examining, and can help refine the experimental design of future studies in order to obtain measurements at time points where a significant change is expected.

6 Discussion

Biomedical longitudinal non-linear data is particularly challenging to analyze due to the likelihood of missing observations and different correlation structures in the data, which limit the use of rm-ANOVA. Although LMEMs have started to replace rm-ANOVA as the choice to analyze biomedical data, both methods yield biased estimates when they are used to fit non-linear data as we have visually demonstrated in Section 3.5. This “model misspecification” error, also known as “Type III” error [17] is particularly important because although the *p-value* is the common measure of statistical significance, the validity of its interpretation is determined by the agreement of the data and the model. Guidelines for statistical reporting in biomedical journals exist (the SAMPL guidelines) [56] but they have not been widely adopted and in the case of longitudinal data, we consider that researchers would benefit from reporting a visual assessment of the correspondence between the model fit and the data, instead of merely relying on a R^2 value.

In this paper we have presented GAMs as a suitable method to analyze non-linear longitudinal data. It is interesting to note that although GAMs are a well established method to analyze temporal data in different fields (among which are palaeoecology, geochemistry, and ecology) [33,52] they are not routinely used in biomedical research despite an early publication from Hastie and Tibshirani that demonstrated their use in medical research [57]. This is possibly due to the fact that the theory behind GAMs can seem very different from that of rm-ANOVA and LMEMs, but the purpose of Section 4 is to demonstrate that at its core the theory quite simple: Instead of using a linear relationship to model the response (as rm-ANOVA and LMEMs do), GAMs use basis functions to build smooths that are capable of following non-linear trends in the data.

However, from a practical standpoint is equally important to demonstrate how GAMs are computationally implemented. We have provided an example on how GAMs can be fitted using simulated data that follows trends reported in biomedical literature [16] using R and the package *mgcv*[37] in Section 5, while a basic workflow for model selection is in the Appendix. One of the features of GAMs is that they go beyond a mere *p-value* to indicate differences between groups, and in turn provide a time-based estimate of shifts in the response that can be directly tied to biological values as the pairwise smooth comparisons in Figure 4 indicate. The model is therefore able to provide an estimate of significant change between the groups at time points where data was not directly measured even with missing data exists (\approx day 4 in Figure 4 A, B), which can be used by researchers as feedback on experiment design and to further evaluate important biological changes in future studies.

We have used R as the software of choice for this paper because not only provides a fully developed environment to fit GAMs, but also eases simulation (which is becoming increasingly used for exploratory statistical analysis and power calculations) and provides powerful and convenient methods of visualization, which are key aspects that biomedical researchers might need to consider to make their work reproducible. In this regard, reproducibility is still an issue in biomedical research [58,59], but it is becoming apparent that what other disciplines have experienced in this aspect is likely to impact soon rather than later this field. Researchers need to plan on how they will make their data, code, and any other materials open and accessible as more journals and funding agencies recognize the importance and benefits of open science in biomedical research. We have made all the data and code used in this paper accessible, and we hope that this will encourage other researchers to do the same with future projects.

7 Conclusion

We have presented GAMs as a method to analyze longitudinal biomedical data. Future directions of this work will include simulation-based estimations of statistical power using GAMs, as well as demonstrating the prediction capabilities of these models using large datasets. By making the data and code used in this paper accessible, we hope to address the need of creating and sharing reproducible work in biomedical research.

8 Acknowledgements

This work was supported by the National Science Foundation Career Award (CBET 1751554) and the Arkansas Biosciences Institute.

9 References

- [1] D. Roblyer, S. Ueda, A. Cerussi, W. Tanamai, A. Durkin, R. Mehta, D. Hsiang, J.A. Butler, C. McLaren, W.P. Chen, B. Tromberg, Optical imaging of breast cancer oxyhemoglobin flare correlates with neoadjuvant chemotherapy response one day after starting treatment, *Proceedings of the National Academy of Sciences of the United States of America*. 108 (2011) 14626–14631. <https://doi.org/10.1073/pnas.1013103108>.
- [2] A. Tank, H.M. Peterson, V. Pera, S. Tabassum, A. Leproux, T. O’Sullivan, E. Jones, H. Cabral, N. Ko, R.S. Mehta, B.J. Tromberg, D. Roblyer, Diffuse optical spectroscopic imaging reveals distinct early breast tumor hemodynamic responses to metronomic and maximum tolerated dose regimens, *Breast Cancer Research*. 22 (2020) 1–10. <https://doi.org/doi:10.1186/s13058-020-01262-1>.
- [3] M.V. Pavlov, T.I. Kalganova, Y.S. Lyubimtseva, V.I. Plekhanov, G.Y. Golubyatnikov, O.Y. Ilyinskaya, A.G. Orlova, P.V. Subochev, D.V. Safonov, N.M. Shakhova, A.V. Maslennikova, Multimodal approach in assessment of the response of breast cancer to neoadjuvant chemotherapy, *Journal of Biomedical Optics*. 23 (2018). <https://doi.org/%7B10.1117/1.JBO.23.9.091410%7D>.
- [4] V. Demidov, A. Maeda, M. Sugita, V. Madge, S. Sadanand, C. Flueraru, I.A. Vitkin, Preclinical longitudinal imaging of tumor microvascular radiobiological response with functional optical coherence tomography, *Scientific Reports*. 8 (2018). <https://doi.org/%7B10.1038/s41598-017-18635-w%7D>.
- [5] G. Ritter, L. Cohen, C. Williams, E. Richards, L. Old, S. Welt, Serological analysis of human anti-human antibody responses in colon cancer patients treated with repeated doses of humanized monoclonal antibody A33, *Cancer Research*. 61 (2001) 6851–6859.
- [6] E.M. Roth, A.C. Goldberg, A.L. Catapano, A. Torri, G.D. Yancopoulos, N. Stahl, A. Brunet, G. Lecorps, H.M. Colhoun, Antidrug antibodies in atients treated with alirocumab, *New England Journal of Medicine*. 376 (2017) 1589–1590. <https://doi.org/%7B10.1056/NEJMc1616623%7D>.
- [7] J.D. Jones, H.E. Ramser, A.E. Woessner, K.P. Quinn, In vivo multiphoton microscopy detects longitudinal metabolic changes associated with delayed skin wound healing, *Communications Biology*. 1 (2018). <https://doi.org/%7B10.1038/s42003-018-0206-4%7D>.
- [8] M.C. Skala, A. Fontanella, L. Lan, J.A. Izatt, M.W. Dewhirst, Longitudinal optical imaging of tumor metabolism and hemodynamics, *Journal of Biomedical Optics*. 15 (2010). <https://doi.org/10.1117/1.3285584>.
- [9] G.J. Greening, K.P. Miller, C.R. Spainhour, M.D. Cato, T.J. Muldoon, Effects of isoflurane anesthesia on physiological parameters in murine subcutaneous tumor allografts measured via diffuse reflectance spectroscopy, *Biomedical Optics Express*. 9 (2018) 2871–2886. <https://doi.org/%7B10.1364/BOE.9.002871%7D>.
- [10] T.T. Sio, P.J. Atherton, B.J. Birkhead, D.J. Schwartz, J.A. Sloan, D.K. Seisler, J.A. Martenson, C.L. Loprinzi, P.C. Griffin, R.F. Morton, J.C. Anders, T.J. Stoffel, R.E. Haselow, R.B. Mowat, M.A.N. Wittich, J.D. Bearden III, R.C. Miller, Repeated measures analyses of dermatitis symptom evolution in breast cancer patients receiving radiotherapy in a phase 3 randomized trial of mometasone furoate vs placebo (N06C4 [alliance]), *Supportive Care in Cancer*. 24 (2016) 3847–3855. <https://doi.org/%7B10.1007/s00520-016-3213-3%7D>.
- [11] J.I. Kamstra, P.U. Dijkstra, M. van Leeuwen, J.L.N. Roodenburg, J.A. Langendijk, Mouth opening in patients irradiated for head and neck cancer: A prospective repeated measures study, *Oral Oncology*. 51 (2015) 548–555. <https://doi.org/%7B10.1016/j.oraloncology.2015.01.016%7D>.
- [12] E.-J. Wagenmakers, M. Lee, T. Lodewyckx, G.J. Iverson, Bayesian versus frequentist inference, in: H. Hoijtink, I. Klugkist, P.A. Boelen (Eds.), *Bayesian Evaluation of Informative Hypotheses*, Springer New York, New York, NY, 2008: pp. 181–207. https://doi.org/10.1007/978-0-387-09612-4_9.

- [13] R. Gueorguieva, J.H. Krystal, Move over ANOVA - Progress in analyzing repeated-measures data and its reflection in papers published in the archives of general psychiatry, *Archives of General Psychiatry*. 61 (2004) 310–317. <https://doi.org/10.1001/archpsyc.61.3.310>.
- [14] P. Schober, T.R. Vetter, Repeated measures designs and analysis of longitudinal data: If at first you do not succeed-try, try again, *Anesthesia and Analgesia*. 127 (2018) 569–575. <https://doi.org/10.1213/ane.0000000000003511>.
- [15] J. Pinheiro, D. Bates, *Mixed-effects models in S and S-PLUS*, Springer Science & Business Media, 2006. <https://doi.org/https://doi.org/10.1007/b98882>.
- [16] K. Vishwanath, H. Yuan, W.T. Barry, M.W. Dewhirst, N. Ramanujam, Using optical spectroscopy to longitudinally monitor physiological changes within solid tumors, *Neoplasia*. 11 (2009) 889–900. <https://doi.org/10.1593/neo.09580>.
- [17] B. Dennis, J.M. Ponciano, M.L. Taper, S.R. Lele, Errors in statistical inference under model misspecification: evidence, hypothesis testing, and AIC, *Frontiers in Ecology and Evolution*. 7 (2019). <https://doi.org/%7B10.3389/fevo.2019.00372%7D>.
- [18] B. Wang, Z. Zhou, H. Wang, X.M. Tu, C. Feng, The p-value and model specification in statistics, *General Psychiatry*. 32 (2019). <https://doi.org/%7B10.1136/gpsych-2019-100081%7D>.
- [19] C. Liu, T.P. Cripe, M.-O. Kim, Statistical issues in longitudinal data analysis for treatment efficacy studies in the biomedical sciences, *Molecular Therapy*. 18 (2010) 1724–1730. <https://doi.org/10.1038/mt.2010.127>.
- [20] L.G. Halsey, D. Curran-Everett, S.L. Vowler, G.B. Drummond, The fickle p value generates irreproducible results, *Nature Methods*. 12 (2015) 179–185. <https://doi.org/%7B10.1038/nmeth.3288%7D>.
- [21] H. Abdi, Holm’s sequential Bonferroni procedure, *Encyclopedia of Research Design*. 1 (2010) 1–8. <https://doi.org/10.4135/9781412961288.n178>.
- [22] S. Nakagawa, A farewell to Bonferroni: the problems of low statistical power and publication bias, *Behavioral Ecology*. 15 (2004) 1044–1045. <https://doi.org/%7B10.1093/beheco/arh107%7D>.
- [23] A. Gelman, J. Hill, M. Yajima, Why we (usually) don’t have to worry about multiple comparisons, *Journal of Research on Educational Effectiveness*. 5 (2012) 189–211. <https://doi.org/%7B10.1080/19345747.2011.618213%7D>.
- [24] C. Albers, The problem with unadjusted multiple and sequential statistical testing, *Nature Communications*. 10 (2019). <https://doi.org/%7B10.1038/s41467-019-09941-0%7D>.
- [25] C. Ugrinowitsch, G.W. Fellingham, M.D. Ricard, Limitations of ordinary least squares models in analyzing repeated measures data, *Medicine and Science in Sports and Exercise*. 36 (2004) 2144–2148. <https://doi.org/10.1249/01.mss.0000147580.40591.75>.
- [26] H. Huynh, L.S. Feldt, Estimation of the Box correction for degrees of freedom from sample data in randomized block and split-Plot designs, *Journal of Educational Statistics*. 1 (1976) 69–82. <https://doi.org/10.3102/10769986001001069>.
- [27] S.W. Greenhouse, S. Geisser, On methods in the analysis of profile data, *Psychometrika*. 24 (1959) 95–112. <https://doi.org/10.1007/bf02289823>.
- [28] N. Haverkamp, A. Beauducel, Violation of the sphericity assumption and its effect on type-I error rates in repeated measures ANOVA and multi-level linear models (MLM), *Frontiers in Psychology*. 8 (2017). <https://doi.org/%7B10.3389/fpsyg.2017.01841%7D>.

- [29] H. Keselman, J. Algina, R. Kowalchuk, The analysis of repeated measures designs: A review, *British Journal of Mathematical & Statistical Psychology*. 54 (2001) 1–20. <https://doi.org/10.1348/0007110011593577D>.
- [30] Jaykaran. Charan, N. Kantharia, How to calculate sample size in animal studies?, *Journal of Pharmacology and Pharmacotherapeutics*. 4 (2013) 303–306. <https://doi.org/10.4103/0976-500X.119726>.
- [31] D.J. Barr, R. Levy, C. Scheepers, H.J. Tily, Random effects structure for confirmatory hypothesis testing: Keep it maximal, *Journal of Memory and Language*. 68 (2013) 255–278. <https://doi.org/10.1016/j.jml.2012.11.0017D>.
- [32] N.L. Rose, H.D. Yang, S.D. Turner, G.L. Simpson, An assessment of the mechanisms for the transfer of lead and mercury from atmospherically contaminated organic soils to lake sediments with particular reference to Scotland, UK, *Geochimica Et Cosmochimica Acta*. 82 (2012) 113–135. <https://doi.org/10.1016/j.gca.2010.12.026>.
- [33] E.J. Pedersen, D.L. Miller, G.L. Simpson, N. Ross, Hierarchical generalized additive models in ecology: An introduction with mgcv, *Peerj*. 7 (2019). <https://doi.org/10.7717/peerj.6876>.
- [34] G.L. Simpson, Modelling palaeoecological time series using generalised additive models, *Frontiers in Ecology and Evolution*. 6 (2018). <https://doi.org/10.3389/fevo.2018.00149>.
- [35] L. Yang, G. Qin, N. Zhao, C. Wang, G. Song, Using a generalized additive model with autoregressive terms to study the effects of daily temperature on mortality, *BMC Medical Research Methodology*. 12 (2012). <https://doi.org/10.1186/1471-2288-12-1657D>.
- [36] N. Beck, S. Jackman, Beyond linearity by default: Generalized additive models, *American Journal of Political Science*. (1998) 596–627.
- [37] S.N. Wood, *Generalized additive models: An introduction with R*, Second Edition, CRC Press LLC, Philadelphia, PA, 2017.
- [38] R Core Team, *R: A language and environment for statistical computing*, R Foundation for Statistical Computing, Vienna, Austria, 2020. <https://www.R-project.org/>.
- [39] S.N. Wood, N. Pya, B. Saeften, Smoothing parameter and model selection for general smooth models, *Journal of the American Statistical Association*. 111 (2016) 1548–1563. <https://doi.org/10.1080/01621459.2016.11809867D>.
- [40] B.T. West, K.B. Welch, A.T. Galecki, *Linear mixed models: A practical guide using statistical software*, second edition, Taylor & Francis, 2014. <https://books.google.com/books?id=hjT6AwAAQBAJ>.
- [41] R.D. Wolfinger, Heterogeneous variance: Covariance structures for repeated measures, *Journal of Agricultural, Biological, and Environmental Statistics*. 1 (1996) 205–230. <http://www.jstor.org/stable/1400366>.
- [42] R.E. Weiss, *Modeling longitudinal data*, Springer New York, 2005. https://books.google.com/books?id=MQ/_bvWDPsEAC.
- [43] S. Geisser, S.W. Greenhouse, An extension of Box’s results on the use of the F distribution in multivariate analysis, *The Annals of Mathematical Statistics*. 29 (1958) 885–891. <https://doi.org/10.1214/aoms/1177706545>.
- [44] S.E. Maxwell, H.D. Delaney, K. Kelley, *Designing experiments and analyzing data: A model comparison perspective*, third edition, Taylor & Francis, 2017. <https://books.google.com/books?id=NmFQDwAAQBAJ>.

- [45] G. Molenberghs, H. Thijs, I. Jansen, C. Beunckens, M. Kenward, C. Mallinckrodt, R. Carroll, Analyzing incomplete longitudinal clinical trial data, *Biostatistics*. 5 (2004) 445–464. <https://doi.org/10.1093/biostatistics/kxh001>.
- [46] J. Scheffer, Dealing with missing data, *Research Letters in the Information and Mathematical Sciences*. 3 (2002) 153–160.
- [47] R.F. Potthoff, G.E. Tudor, K.S. Pieper, V. Hasselblad, Can one assess whether missing data are missing at random in medical studies?, *Statistical Methods in Medical Research*. 15 (2006) 213–234. <https://doi.org/10.1191/0962280206sm448oa>.
- [48] Y. Ma, M. Mazumdar, S.G. Memtsoudis, Beyond repeated-measures analysis of variance advanced statistical methods for the analysis of longitudinal data in anesthesia research, *Regional Anesthesia and Pain Medicine*. 37 (2012) 99–105. <https://doi.org/10.1097/AAP.0b013e31823ebc74>.
- [49] J. Pinheiro, D. Bates, S. DebRoy, D. Sarkar, R Core Team, nlme: Linear and nonlinear mixed effects models, 2020. <https://CRAN.R-project.org/package=nlme>.
- [50] J.A. Nelder, R.W.M. Wedderburn, Generalized linear models, *Journal of the Royal Statistical Society. Series A (General)*. 135 (1972) 370–384. <http://www.jstor.org/stable/2344614>.
- [51] T. Hastie, R. Tibshirani, Generalized additive models: Some applications, *Journal of the American Statistical Association*. 82 (1987) 371–386. <https://doi.org/10.1080/01621459.1987.10478440>.
- [52] T.J. Hefley, K.M. Broms, B.M. Brost, F.E. Buderman, S.L. Kay, H.R. Scharf, J.R. Tipton, P.J. Williams, M.B. Hooten, The basis function approach for modeling autocorrelation in ecological data, *Ecology*. 98 (2017) 632–646. <https://doi.org/10.1002/ecy.1674>.
- [53] E.J. Wegman, I.W. Wright, Splines in statistics, *Journal of the American Statistical Association*. 78 (1983) 351–365. <https://doi.org/10.1080/01621459.1983.10477977>.
- [54] G.L. Simpson, Gratia: Graceful 'ggplot'-based graphics and other functions for GAMs fitted using 'mgcv', 2020. <https://CRAN.R-project.org/package=gratia>.
- [55] J. Harezlak, D. Ruppert, M.P. Wand, *Semiparametric Regression with R*, Springer New York, 2018. <https://doi.org/10.1007/978-1-4939-8853-2>.
- [56] T.A. Lang, D.G. Altman, Basic statistical reporting for articles published in Biomedical Journals: The “Statistical Analyses and Methods in the Published Literature” or the SAMPL Guidelines, *INTERNATIONAL JOURNAL OF NURSING STUDIES*. 52 (2015) 5–9. <https://doi.org/10.1016/j.ijnurstu.2014.09.006>.
- [57] T. Hastie, R. Tibshirani, Generalized additive models for medical research, *Statistical Methods in Medical Research*. 4 (1995) 187–196. <https://doi.org/10.1177/096228029500400302>.
- [58] C.G. Begley, J.P.A. Ioannidis, Reproducibility in Science Improving the Standard for Basic and Preclinical Research, *Circulation Research*. 116 (2015) 116–126. <https://doi.org/10.1161/CIRCRESAHA.114.303819>.
- [59] T.L. Weissgerber, O. Garcia-Valencia, V.D. Garovic, N.M. Milic, S.J. Winham, Meta-Research: Why we need to report more than 'Data were Analyzed by t-tests or ANOVA', *Elife*. 7 (2018) e36163. <https://doi.org/10.7554/eLife.36163>.

A Code for Manuscript data

This section presents the code used to generate figures, models and simulated data from Sections 3 and 4 from the main manuscript.

A.1 Compound symmetry and independent errors in linear and quadratic responses

This section simulated linear and quadratic data in the same manner as in Section 3.5. The linear simulations using Figure 5 show in panels A and D the simulated mean responses and individual data points. Panels C and G show a visual interpretation of “correlation” in the responses: In panel C, subjects that have a value of the random error ε either above or below the mean group response are more likely to have other observations that follow the same trajectory, thereby demonstrating correlation in the response. In panel G, because the errors are independent, there is no expectation that responses are likely to follow a similar pattern. Panels D and H show the predictions from the rm-ANOVA model.

The following code produces a more comprehensive exploration of Figure 1 in the main manuscript.

```
set.seed(1)
#####Section for calculations#####

## Example with linear response

#This function simulates data using a linear or quadratic mean response
  and each with correlated
#or uncorrelated errors. Each group has a different slope/concavity.
example <- function(n_time = 6, #number of time points
                    fun_type = "linear", #type of response
                    error_type = "correlated") {

  if (!(fun_type %in% c("linear", "quadratic")))
    stop('fun_type must be either "linear", or "quadratic"')
  if (!(error_type %in% c("correlated", "independent")))
    stop('fun_type must be either "correlated", or "independent"')

  x <- seq(1,6, length.out = n_time)

  #Create mean response matrix: linear or quadratic
  mu <- matrix(0, length(x), 2)
  # linear response
  if (fun_type == "linear") {
    mu[, 1] <- - (0.25*x)+2
    mu[, 2] <- 0.25*x+2
  } else {
    # quadratic response (non-linear)

    mu[, 1] <- -(0.25 * x^2) +1.5*x-1.25
    mu[, 2] <- (0.25 * x^2) -1.5*x+1.25
  }

  #create an array where individual observations per each time point for
    each group are to be stored. Currently using 10 observations per
    timepoint
  y <- array(0, dim = c(length(x), 2, 10))

  #Create array to store the "errors" for each group at each timepoint.
    The "errors" are the
  #between-group variability in the response.
```

```

674 errors <- array(0, dim = c(length(x), 2, 10))
675 #create an array where 10 observations per each time point for each
676     group are to be stored
677
678 #The following cycles create independent or correlated responses. To
679     each value of mu (mean response per group) a randomly generated error
680     (correlated or uncorrelated) is added and thus the individual
681     response is created.
682 if (error_type == "independent") {
683     ## independent errors
684     for (i in 1:2) {
685         for (j in 1:10) {
686             errors[, i, j] <- rnorm(6, 0, 0.25)
687             y[, i, j] <- mu[, i] + errors[, i, j]
688         }
689     }
690 } else {
691     for (i in 1:2) {      # number of treatments
692         for (j in 1:10) { # number of subjects
693             # compound symmetry errors: variance covariance matrix
694             errors[, i, j] <- rmvn(1, rep(0, length(x)), 0.1 * diag(6) + 0.25
695                 * matrix(1, 6, 6))
696             y[, i, j] <- mu[, i] + errors[, i, j]
697         }
698     }
699 }
700
701
702 ## subject random effects
703
704 ## visualizing the difference between independent errors and compound
705     symmetry
706 ## why do we need to account for this -- overly confident inference
707
708 #labeling y and errors
709 dimnames(y) <- list(time = x,
710                     treatment = 1:2,
711                     subject = 1:10)
712
713 dimnames(errors) <- list(time = x,
714                          treatment = 1:2,
715                          subject = 1:10)
716
717 #labeling the mean response
718 dimnames(mu) <- list(time = x,
719                     treatment = 1:2)
720
721 #convert y, mu and errors to dataframes with time, treatment and
722     subject columns
723 dat <- as.data.frame.table(y,
724                           responseName = "y")
725 dat_errors <- as.data.frame.table(errors,
726                                  responseName = "errors")
727 dat_mu <- as.data.frame.table(mu,

```

```

728         responseName = "mu")
729
730 #join the dataframes to show mean response and errors per subject
731 dat <- left_join(dat, dat_errors,
732                 by = c("time", "treatment", "subject"))
733 dat <- left_join(dat, dat_mu,
734                 by = c("time", "treatment"))
735 #add time
736 dat$time <- as.numeric(as.character(dat$time))
737 #label subjects per group
738 dat <- dat %>%
739     mutate(subject = factor(paste(subject,
740                                   treatment,
741                                   sep = "-")))
742
743
744 ## repeated measures ANOVA
745
746 fit_anova <- lm(y ~ time + treatment + time * treatment, data = dat)
747
748 #LMEM: time and treatment interaction model, compound symmetry
749 fit_lme <- lme(y ~ treatment + time + treatment:time,
750               data = dat,
751               random = ~ 1 | subject,
752               correlation = corCompSymm(form = ~ 1 | subject)
753             )
754
755 #create a prediction frame where the model can be used for plotting
756 #purposes
757 pred_dat <- expand.grid(
758   treatment = factor(1:2),
759   time = unique(dat$time)
760 )
761
762 #add model predictions to the dataframe that has the simulated data
763 dat$pred_anova <- predict(fit_anova)
764 dat$pred_lmem <- predict(fit_lme)
765
766 #return everything in a list
767 return(list(
768   dat = dat,
769   pred_dat = pred_dat,
770   fit_anova=fit_anova,
771   fit_lme = fit_lme
772 ))
773 }
774 #####Section for plotting#####
775 #####
776 #This function will create the plots for either a "linear" or "quadratic"
777 #response
778
779 plot_example <- function(sim_dat) {
780   ## Plot the simulated data (scatterplot)
781

```

```

782 p1 <- sim_dat$dat %>%
783   ggplot(aes(x = time,
784             y = y,
785             group = treatment,
786             color = treatment)
787         ) +
788   geom_point(show.legend=FALSE) +
789   labs(y='response')+
790   geom_line(aes(x = time,
791               y = mu,
792               color = treatment),
793             show.legend=FALSE) +
794   theme_classic() +
795   theme(plot.title = element_text(size = 30,
796                                   face = "bold"),
797         text=element_text(size=30))+
798   thm
799
800 #plot the simulated data with trajectories per each subject
801 p2 <- sim_dat$dat %>%
802   ggplot(aes(x = time,
803             y = y,
804             group = subject,
805             color = treatment)
806         ) +
807   geom_line(aes(size = "Subjects"),
808             show.legend = FALSE) +
809   # facet_wrap(~ treatment) +
810   geom_line(aes(x = time,
811               y = mu,
812               color = treatment,
813               size = "Simulated Truth"),
814             lty = 1, show.legend = FALSE) +
815   labs(y='response')+
816   scale_size_manual(name = "Type", values=c("Subjects" = 0.5, "Simulated
817       Truth" = 3)) +
818   theme_classic()+
819   theme(plot.title = element_text(size = 30,
820                                   face = "bold"),
821         text=element_text(size=30))+
822   thm
823
824 #plot the errors
825 p3 <- sim_dat$dat %>%
826   ggplot(aes(x = time,
827             y = errors,
828             group = subject,
829             color = treatment)) +
830   geom_line(show.legend=FALSE) +
831   labs(y='errors')+
832   theme_classic()+
833   theme(plot.title = element_text(size = 30,
834                                   face = "bold"),
835         text=element_text(size=30))+

```



```

836   thm
837
838   #plot the model predictions for rm-ANOVA
839   p4 <- ggplot(sim_dat$dat,
840               aes(x = time,
841                   y = y,
842                   color = treatment)) +
843   geom_point(show.legend=FALSE)+
844   labs(y='response')+
845   geom_line(aes(y = predict(sim_dat$fit_anova),
846                   group = subject, size = "Subjects"),show.legend = FALSE)
847   +
848   geom_line(data = sim_dat$pred_dat,
849             aes(y = predict(sim_dat$fit_anova,
850                             level = 0,
851                             newdata = sim_dat$pred_dat),
852                 size = "Population"),
853             show.legend=FALSE) +
854   guides(color = guide_legend(override.aes = list(size = 2)))+
855   scale_size_manual(name = "Predictions",
856                     values=c("Subjects" = 0.5, "Population" = 3)) +
857   theme_classic() +
858   theme(plot.title = element_text(size = 30,
859                                   face = "bold"),
860         text=element_text(size=30))+
861   thm
862
863
864
865   #plot the LMEM predictions
866   p5 <- ggplot(sim_dat$dat,
867               aes(x = time,
868                   y = y,
869                   color = treatment)) +
870   geom_point()+
871   labs(y='response')+
872   geom_line(aes(y = predict(sim_dat$fit_lme),
873                   group = subject, size = "Subjects")) +
874   geom_line(data = sim_dat$pred_dat,
875             aes(y = predict(sim_dat$fit_lme,
876                             level = 0,
877                             newdata = sim_dat$pred_dat),
878                 size = "Population")) +
879   guides(color = guide_legend(override.aes = list(size = 2)))+
880   scale_size_manual(name = "Predictions",
881                     values=c("Subjects" = 0.5, "Population" = 3)) +
882   theme_classic() +
883   theme(plot.title = element_text(size = 30,
884                                   face = "bold"),
885         text=element_text(size=30))+
886   thm
887
888   return((p1+p3+p2+p4+p5)+plot_layout(nrow=1)+plot_annotation(tag_levels =
889   'A'))

```

```

890
891
892 }
893
894 txt<-18
895
896 #Store each plot in a separate object
897 A1<-plot_example(example(fun_type = "linear", error_type = "correlated"))
898
899 B1<-plot_example(example(fun_type = "linear", error_type = "independent"))
900
901 C1<-plot_example(example(fun_type = "quadratic", error_type = "correlated"
902   ))
903
904 D1<-plot_example(example(fun_type = "quadratic", error_type = "independent
905   "))
906

```

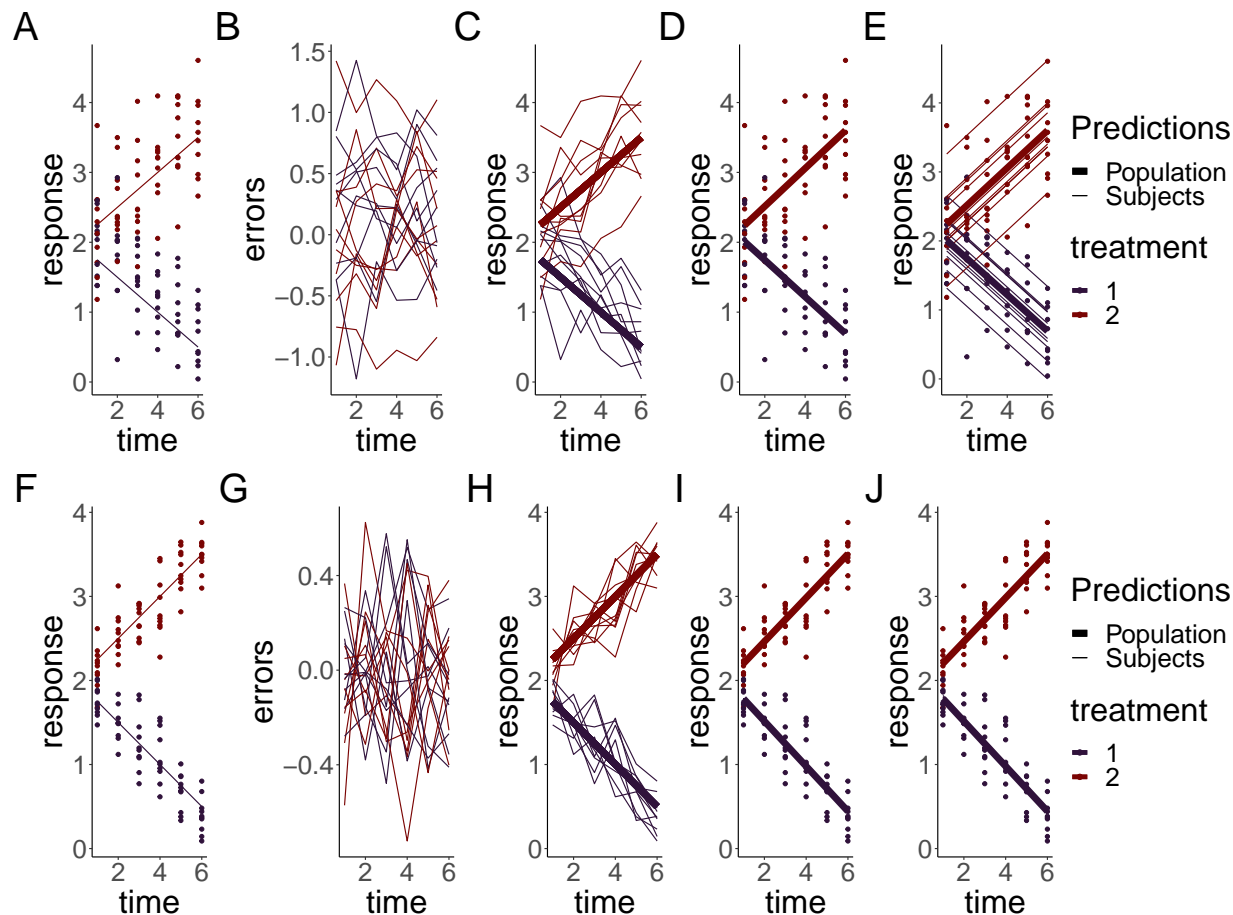


Figure 5: Simulated linear responses from two groups with correlated (top row) or independent (bottom row) errors using a rm-ANOVA model and a LMEM. A, F: Simulated data with known mean response and individual responses (points) showing the dispersion of the data. B, G: Generated errors showing the difference in the behavior of correlated and independent errors. C, H: Simulated data with thin lines representing individual trajectories. D, I: Estimations from the rm-ANOVA model for the mean group response. E, J: Estimations from the LMEM for the mean group response and individual responses (thin lines). In all panels, thick lines are the predicted mean response per group, thin lines are the random effects for each subject and points represent the original raw data. Both rm-ANOVA and the LMEM are able to capture the trend of the data.

907 For the quadratic response case, Figure 6 shows the simulated responses using compound symmetry and
 908 independent errors.

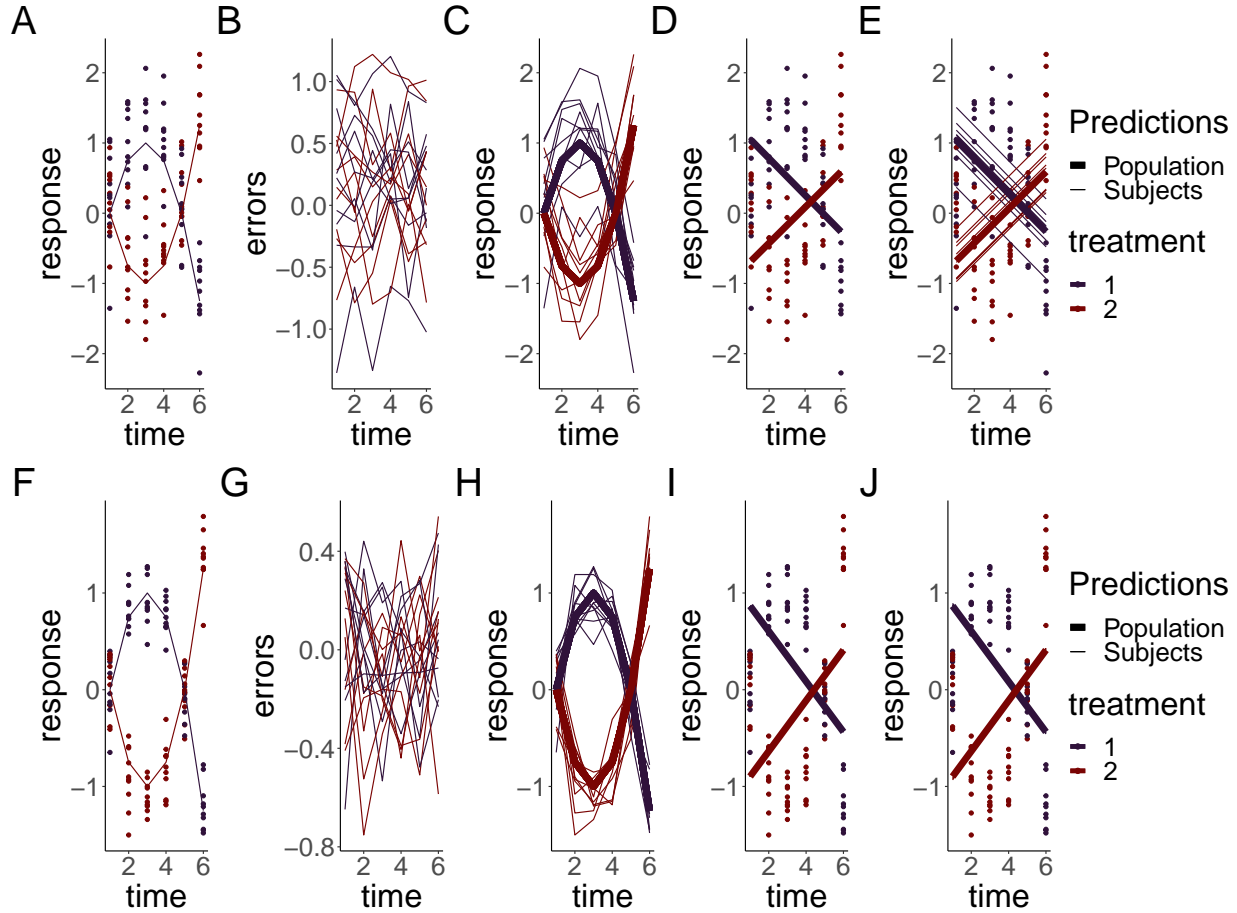


Figure 6: Simulated quadratic responses from two groups with correlated (top row) or independent (bottom row) errors using a rm-ANOVA model and a LMEM. A, F: Simulated data with known mean response and individual responses (points) showing the dispersion of the data. B, G: Generated errors showing the difference in the behavior of correlated and independent errors. C, H: Simulated data with thin lines representing individual trajectories. D, I: Estimations from the rm-ANOVA model for the mean group response. E, J: Estimations from the LMEM for the mean group response and individual responses (thin lines). In all panels, thick lines are the predicted mean response per group, thin lines are the random effects for each subject and points represent the original raw data. Both rm-ANOVA and the LMEM are not able to capture the changes in each group over time.

A.2 Basis functions and GAMs

This code produces Figure 2 from the main manuscript. Briefly, a non-linear (quadratic) response is simulated a gam model is fitted and the basis are extracted in order to explain how the smooth is constructed. The code for data simulation is used again here for the sake of keeping the same structure, although the data can be simulated in a more simple fashion.

```
#generate the response: the same initial procedure from the previous
#section to simulate
#the response
set.seed(1)
n_time = 6
x <- seq(1,6, length.out = n_time)
mu <- matrix(0, length(x), 2)
mu[, 1] <- -(0.25 * x^2) + 1.5*x - 1.25 #mean response
```

```

923 mu[, 2] <- (0.25 * x^2) -1.5*x+1.25 #mean response
924 y <- array(0, dim = c(length(x), 2, 10))
925 errors <- array(0, dim = c(length(x), 2, 10))
926 for (i in 1:2) { # number of treatments
927   for (j in 1:10) { # number of subjects
928     # compound symmetry errors
929     errors[, i, j] <- rmvn(1, rep(0, length(x)), 0.1 * diag(6) + 0.25
930       * matrix(1, 6, 6))
931     y[, i, j] <- mu[, i] + errors[, i, j]
932   }
933 }
934
935 #label each table
936 dimnames(y) <- list(time = x, treatment = 1:2, subject = 1:10)
937 dimnames(errors) <- list(time = x, treatment = 1:2, subject = 1:10)
938 dimnames(mu) <- list(time = x, treatment = 1:2)
939
940 #Convert to dataframes with subject, time and group columns
941 dat <- as.data.frame.table(y, responseName = "y")
942 dat_errors <- as.data.frame.table(errors, responseName = "errors")
943 dat_mu <- as.data.frame.table(mu, responseName = "mu")
944 dat <- left_join(dat, dat_errors, by = c("time", "treatment", "subject"))
945 dat <- left_join(dat, dat_mu, by = c("time", "treatment"))
946 dat$time <- as.numeric(as.character(dat$time))
947
948 #label subject per group
949 dat <- dat %>%
950   mutate(subject = factor(paste(subject, treatment, sep = "-")))
951
952 #extract "Group 1" to fit the GAM
953 dat<-subset(dat,treatment==1)
954 #keep just the response and timepoint columns
955 dat<-dat[,c('y','time')]
956
957 #GAM model of time, 5 knots
958 gm<-gam(y~s(time,k=5),data=dat)
959
960 #model_matrix (also known as) 'design matrix'
961 #will contain the smooths used to create model 'gm'
962 model_matrix<-as.data.frame(predict(gm,type='lpmatrix'))
963
964
965 time<-c(1:6)
966
967 basis<-model_matrix[1:6,] #extracting basis (because the values are
968   repeated after every 6 rows)
969 #basis<-model_matrix[1:6,-1] #extracting basis
970 colnames(basis)[colnames(basis)=="(Intercept)"]<- "s(time).0"
971 basis<-basis %>% #pivoting to long format
972   pivot_longer(
973     cols=starts_with("s")
974   )%>%
975   arrange(name) #ordering
976

```

```

977 #length of dataframe to be created: number of knots by number of
978     timepoints (minus 1 for the intercept that we won't plot)
979 ln<-6*(length(coef(gm)))
980
981 basis_plot<-data.frame(Basis=integer(ln),
982                        value_orig=double(ln),
983                        time=integer(ln),
984                        cof=double(ln)
985 )
986
987 basis_plot$time<-rep(time) #pasting timepoints
988 basis_plot$Basis<-factor(rep(c(1:5),each=6)) #pasting basis number values
989 basis_plot$value_orig<-basis$value #pasting basis values
990 basis_plot$cof<-rep(coef(gm)[1:5],each=6) #pasting coefficients
991 basis_plot<-basis_plot%>%
992     mutate(mod_val=value_orig*cof) #the create the predicted values the
993     bases need to be
994 #multiplied by the coefficients
995
996 #creating labeller to change the labels in the basis plots
997
998 basis_names<-c(
999     '1'="Intercept",
1000     '2'="1",
1001     '3'="2",
1002     '4'="3",
1003     '5'="4"
1004 )
1005
1006 #calculating the final smooth by aggregating the basis functions
1007
1008 smooth<-basis_plot%>%
1009     group_by(time)%>%
1010     summarize(smooth=sum(mod_val))
1011
1012
1013 #original basis
1014 sz<-1
1015 p11<-ggplot(basis_plot,
1016             aes(x=time,
1017                 y=value_orig,
1018                 colour=as.factor(Basis)
1019             )
1020             )+
1021     geom_line(size=sz,
1022              show.legend=FALSE)+
1023     geom_point(size=sz+1,
1024               show.legend = FALSE)+
1025     labs(y='Basis functions')+
1026     facet_wrap(~Basis,
1027               labeller = as_labeller(basis_names)
1028             )+
1029     theme_classic()+
1030     thm

```

```

1031
1032
1033 #penalized basis
1034 p12<-ggplot(basis_plot,
1035             aes(x=time,
1036                 y=mod_val,
1037                 colour=as.factor(Basis)
1038             )
1039         )+
1040     geom_line(show.legend = FALSE,
1041              size=sz)+
1042     geom_point(show.legend = FALSE,
1043               size=sz+1)+
1044     labs(y='Penalized \n basis functions')+
1045     scale_y_continuous(breaks=seq(-1,1,1))+
1046     facet_wrap(~Basis,
1047               labeller=as_labeller(basis_names)
1048             )+
1049     theme_classic()+
1050     thm
1051
1052 #heatmap of the coefficients
1053 x_labels<-c("Intercept","1","2","3","4")
1054 p13<-ggplot(basis_plot,
1055             aes(x=Basis,
1056                 y=Basis))+
1057     geom_tile(aes(fill = cof),
1058              colour = "black") +
1059     scale_fill_gradient(low = "white",
1060                        high = "#B50A2AFF")+ #color picked from KikiMedium
1061     labs(x='Basis',
1062          y='Basis')+
1063     scale_x_discrete(labels=x_labels)+
1064     geom_text(aes(label=round(cof,2)),
1065              size=7,
1066              show.legend = FALSE)+
1067     theme_classic()+
1068     theme(legend.title = element_blank())
1069
1070 #plotting simulated datapoints and smooth term
1071 p14<-ggplot(data=dat,
1072             aes(x=time,y=y))+
1073     geom_point(size=sz+1)+
1074     labs(y='Simulated \n response')+
1075     geom_line(data=smooth,
1076              aes(x=time,
1077                  y=smooth),
1078              color="#6C581DFF",
1079              size=sz+1)+
1080     theme_classic()
1081
1082
1083 #Combining all
1084 b_plot<-p11+p13+p12+p14+plot_annotation(tag_levels='A')&

```

```

1085 theme(
1086   text=element_text(size=18)
1087 )
1088

```

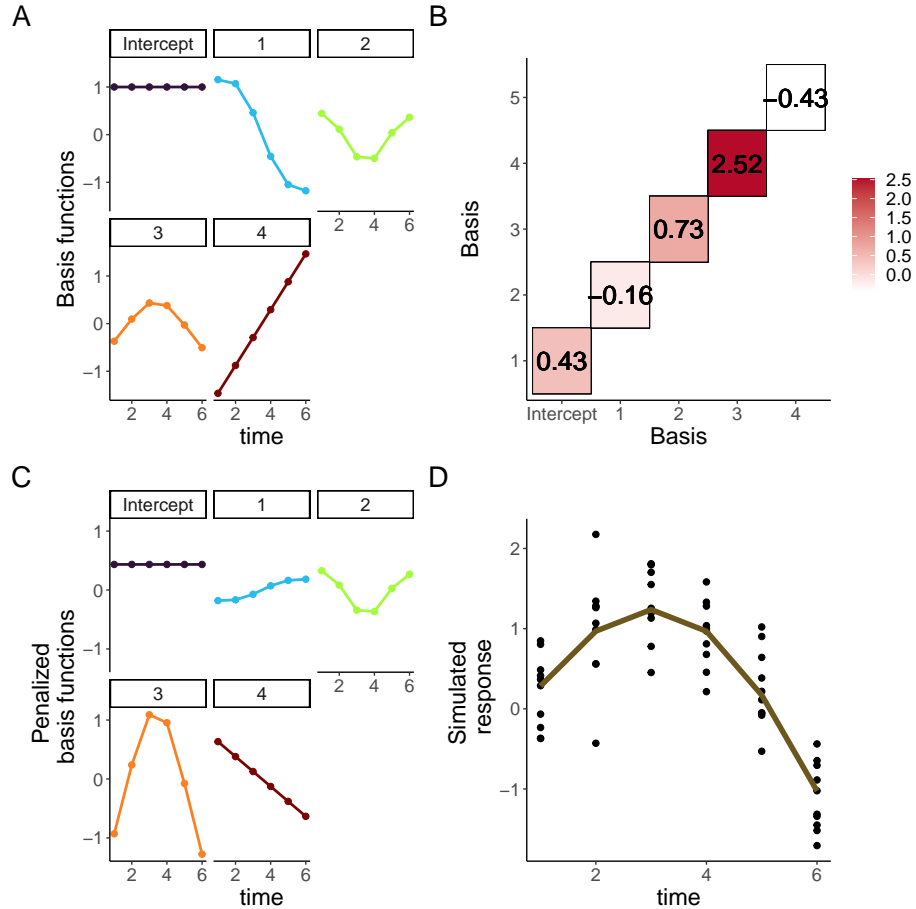


Figure 7: Basis functions for a single smoother for time with five knots. A: Basis functions for a single smoother for time for the simulated data of Group 1 from Figure 2, the intercept basis is not shown. B: Penalty matrix for the basis functions. Each basis function is penalized by a coefficient which can be positive or negative. The coefficient determines the overall effect of each basis in the final smoother. C: Penalized basis functions. Each of the four basis functions of panel A has been penalized by the corresponding coefficient shown in Panel B, note the corresponding increase (or decrease) of each basis. D: Smoother for time and original data points. The smoother (line) is the result of the sum of each penalized basis function at each time point, simulated values for the group appear as points.

B Longitudinal biomedical data simulation and GAMs

This section describes how to fit GAMs to longitudinal data using simulated data. First, data is simulated according to Section 5, where reported data of oxygen saturation (StO_2) in tumors under either chemotherapy or saline control is used as a starting point to generate individual responses in each group.

```

1093 set.seed(1)
1094 #Dataframe that contains the original reported trends
1095 dat<-tibble(StO2=c(4,27,3,2,0.5,7,4,50,45,56),
1096             Day=rep(c(0,2,5,7,10),times=2),

```



```

1098         Group=as.factor(rep(c("Control","Treatment"),each=5))
1099     )
1100
1101
1102     ## plot the mean response
1103     f1<-ggplot(dat,
1104               aes(x = Day,
1105                  y = St02,
1106                  color = Group)) +
1107       geom_line(size=1,
1108                show.legend = FALSE)+
1109       geom_point(show.legend = FALSE,
1110                 size=1.5,
1111                 alpha=0.5)+
1112       labs(y=expression(paste(St0[2],
1113                               ' (real)')))+
1114       theme_classic()+
1115       thm+
1116       scale_x_continuous(breaks=c(0,5,10))+
1117       scale_y_continuous(breaks=c(0,40))+
1118       plot_layout(tag_level = 'new')+
1119       theme(
1120         plot.background = element_rect(fill = "transparent",
1121                                         color = NA),
1122         axis.text=element_text(size=14)
1123       )
1124
1125
1126     #This function simulates data for the tumor data using default parameters
1127     #of 10 observations per time point, and Standard deviation (sd) of 5%.
1128     #Because physiologically St02 cannot go below 0%, data is generated with
1129     #a cutoff value of 0.0001 (the "St02_sim")
1130
1131     simulate_data <- function(dat, n = 10, sd = 5) {
1132       dat_sim <- dat %>%
1133         slice(rep(1:n(), each = n)) %>%
1134         group_by(Group, Day) %>%
1135         mutate(
1136           St02_sim = pmax(rnorm(n, St02, sd), 0.0001),
1137           subject=rep(1:10),
1138           subject=factor(paste(subject, Group, sep = "-"))
1139         ) %>%
1140         ungroup()
1141
1142       return(dat_sim)
1143     }
1144
1145
1146     #subject = factor(paste(subject, treatment, sep = "-"))
1147     n <- 10 #number of observations
1148     sd <- 10 #approximate sd from paper
1149     df <- 6
1150     dat_sim <- simulate_data(dat, n, sd)
1151

```

```

1152 #plotting simulated data
1153 f2<-ggplot(dat_sim,
1154           aes(x = Day,
1155               y = StO2_sim,
1156               color = Group)) +
1157   geom_point(show.legend=FALSE,
1158             size=1.5,
1159             alpha=0.5)+
1160   stat_summary(aes(y = StO2_sim,
1161                   group=Group),
1162               fun=mean, geom="line",
1163               size=1,
1164               show.legend = FALSE)+
1165   labs(y=expression(atop(StO2 ,
1166                          '(simulated)')))+
1167   theme_classic()+
1168   theme(
1169     axis.text=element_text(size=22)
1170   )+
1171   thm+
1172   scale_x_continuous(breaks=c(0,2,5,7,10))
1173

```

1174 B.1 A basic Workflow for GAMs

1175 This section shows a basic workflow to fit a series of increasingly complex GAMs to the simulated data from
 1176 the previous section. Graphical and parameter diagnostics for goodness of fit are discussed, as well as model
 1177 comparison via AIC (Aikake Information Criterion).

1178 B.1.1 First model

1179 The first model fitted to the data is one that only accounts for different smooths by day. The model syntax
 1180 specifies that `gam_00` is the object that will contain all the model information, and that the model attempts
 1181 to explain changes in `StO2_sim` (simulated StO₂) using a smooth per `Day`. The model will use 5 knots (`k=5`)
 1182 for the smooth. The smooth is constructed by default using thin plate regression splines. The smoothing
 1183 parameter estimation method used is the restricted maximum likelihood (REML).

```

1184
1185 gam_00<-gam(StO2_sim ~ s(Day, k = 5),
1186            method='REML',
1187            data = dat_sim)
1188

```

1189 To obtain model diagnostics, two methodologies are to be used: 1) graphical diagnostics, and 2) a model
 1190 check. In the first case, the functions `appraise` and `draw` from the package *gratia* can be used to obtain
 1191 a single output with all the graphical diagnostics. For model check, the functions `gam.check` and `summary`
 1192 from *mgcv* provide detailed information about the model fit and its parameters.

```

1193
1194 #see https://patchwork.data-imaginist.com/reference/area.html
1195 layout1 <- c(
1196   area(1, 1),
1197   area( 1, 2),
1198   area(2, 1),
1199   area(2, 2),
1200   area(1, 3, 2)

```

```

1201 )
1202
1203
1204 layout2 <- c(
1205   area(1, 1),
1206   area( 1, 2),
1207   area(2, 1),
1208   area(2, 2),
1209   area(1,3,2,5)
1210 )
1211
1212 #plot(layout2)
1213

```

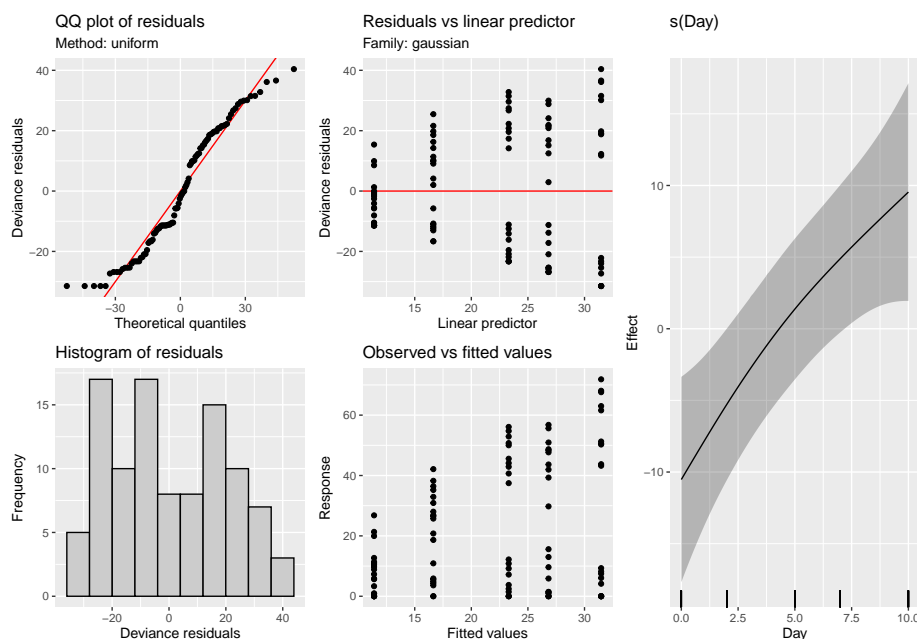


Figure 8: Graphical diagnostics for the first GAM model. Left: Graphical diagnostics provided by the function `appraise` from the package *gratia*. Right: Fitted smooth for the model, provided by the function `draw`.

B.1.1.1 Graphical diagnostics From the output of the function `appraise` in Figure 8, the major indicators of concern about the model are the QQ plot of residuals and the histogram of residuals. The QQ plot shows that the errors are not reasonably located along the 45° line (which indicates normality), as there are multiple points that deviate from the trend, specially in the tails. The histogram also shows that the variation (residuals) is not following the assumption of a normal distribution.

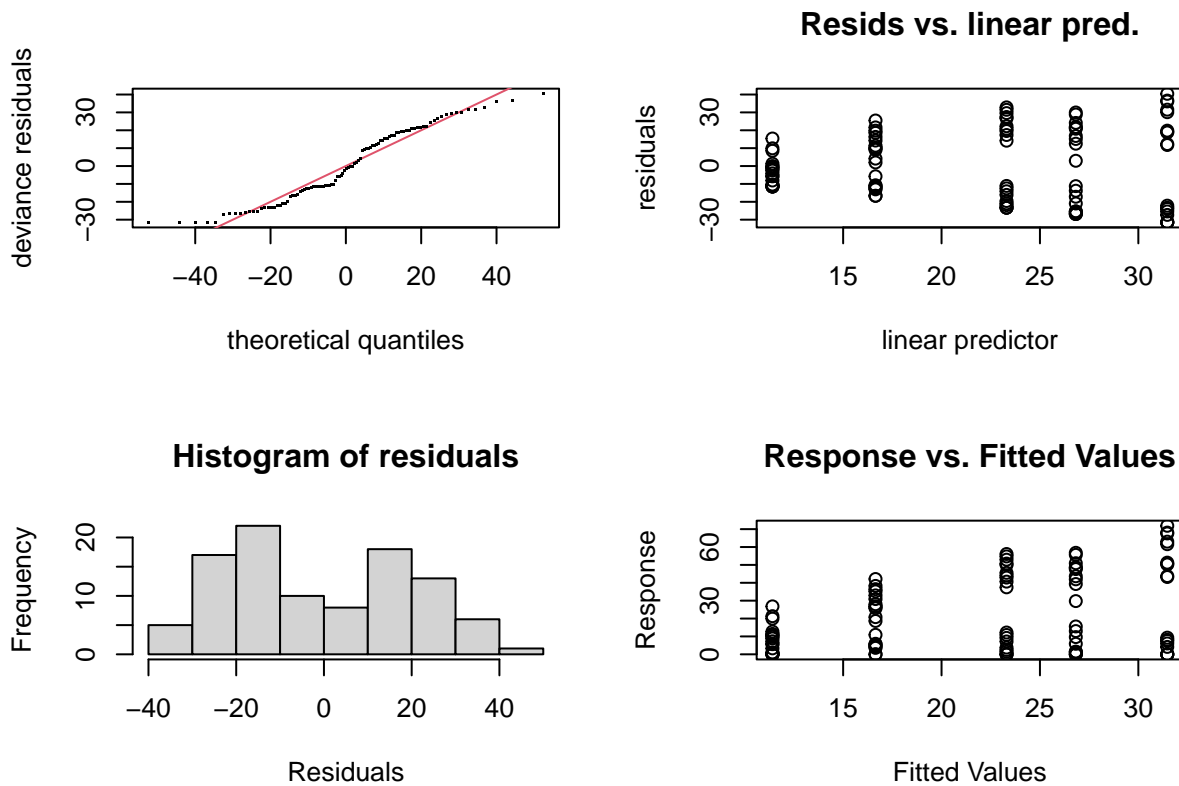
The `draw` function permits to plot the smooths as `ggplot2` objects, which eases subsequent manipulation, if desired. Because model `gam_00` specifies only one smooth for the time covariate (Day), the plot only contains only one smooth. Note that the smooth shows an almost linear profile.

B.1.1.2 Model check

```

1222 #need to add figure number and caption
1223 gam.check(gam_00)
1224

```



```

1226
1227
1228 ##
1229 ## Method: REML   Optimizer: outer newton
1230 ## full convergence after 5 iterations.
1231 ## Gradient range [-6.11034e-07,-1.169842e-07]
1232 ## (score 439.1428 & scale 414.047).
1233 ## Hessian positive definite, eigenvalue range [0.05006795,49.00066].
1234 ## Model rank = 5 / 5
1235 ##
1236 ## Basis dimension (k) checking results. Low p-value (k-index<1) may
1237 ## indicate that k is too low, especially if edf is close to k'.
1238 ##
1239 ##      k'   edf k-index p-value
1240 ## s(Day) 4.00 1.36    0.26 <2e-16 ***
1241 ## ---
1242 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1243

```

```

1244
1245 summary(gam_00)
1246

```

```

1247
1248 ##
1249 ## Family: gaussian
1250 ## Link function: identity
1251 ##
1252 ## Formula:
1253 ## StO2_sim ~ s(Day, k = 5)
1254 ##
1255 ## Parametric coefficients:

```

```

1256 ##           Estimate Std. Error t value Pr(>|t|)
1257 ## (Intercept)    21.929      2.035   10.78  <2e-16 ***
1258 ## ---
1259 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1260 ##
1261 ## Approximate significance of smooth terms:
1262 ##           edf Ref.df      F p-value
1263 ## s(Day)  1.359   1.624  6.695 0.00273 **
1264 ## ---
1265 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1266 ##
1267 ## R-sq.(adj) =  0.106   Deviance explained = 11.8%
1268 ## -REML = 439.14   Scale est. = 414.05      n = 100
1269

```

1270 Special attention must be paid to the ‘k-index’ from `gam.check`. This parameter indicates if the basis
1271 dimension of the smooth is adequate, i.e., it checks that the basis used to create the smooth are adequate
1272 to capture the trends in the data. If the model is not adequately capturing the trends in the data, this is
1273 indicated by a low k-index value (<1). From the output, it can be seen that the k-index is effectively <0.3,
1274 which indicates that the model is not capturing the variability in the data. The ‘edf’ (effective degrees of
1275 freedom) is an indicator of the complexity of the smooth. Here the complexity of the smooth is comparable
1276 to that of a 4th degree polynomial.

1277 From the `summary` function, information about the assumed distribution of the errors (Gaussian in this
1278 case) and the link function can be obtained. The link function is ‘identity’ as the model does not make
1279 any transformation on the predictors. The ‘significance of smooth terms’ *p-value* indicates if each smooth
1280 is adding significance to the model. Here, the *p-value* is low but we have seen that there are issues with
1281 the model from the previous outputs. Finally, the ‘deviance explained’ indicates how much of the data the
1282 model is able to capture, which in this case corresponds to $\approx 12\%$.

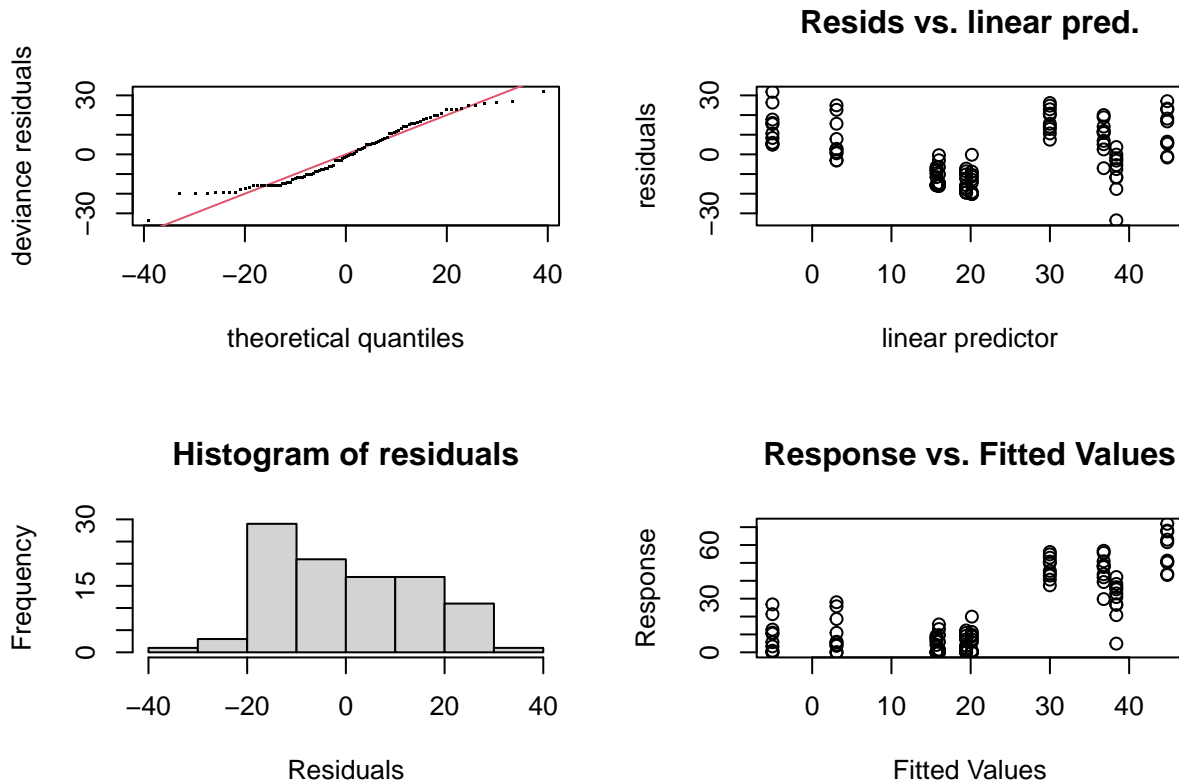
1283 B.1.2 Second model

1284 The major flaw of `gam_00` is that this model is not taking into account the fact that the data is nested in
1285 groups. The next iteration is a model where a different smooth of time (Day) is assigned for each group
1286 using `by=Group` in the model syntax.

```

1287 gam_01<-gam(St02_sim ~ s(Day, by=Group,k = 5),
1288             method='REML',
1289             data = dat_sim)
1290
1291 gam.check(gam_01)
1292
1293

```



```

1294
1295 ##
1296 ##
1297 ## Method: REML   Optimizer: outer newton
1298 ## full convergence after 11 iterations.
1299 ## Gradient range [-3.57434e-06,1.383186e-06]
1300 ## (score 413.523 & scale 230.4732).
1301 ## Hessian positive definite, eigenvalue range [0.1532335,48.55232].
1302 ## Model rank = 9 / 9
1303 ##
1304 ## Basis dimension (k) checking results. Low p-value (k-index<1) may
1305 ## indicate that k is too low, especially if edf is close to k'.
1306 ##
1307 ##           k'   edf k-index p-value
1308 ## s(Day):GroupControl  4.00 3.49    0.38 <2e-16 ***
1309 ## s(Day):GroupTreatment 4.00 2.96    0.38 <2e-16 ***
1310 ## ---
1311 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1312

```

```

1313
1314 summary(gam_01)
1315
1316 ##
1317 ##
1318 ## Family: gaussian
1319 ## Link function: identity
1320 ##
1321 ## Formula:
1322 ## StO2_sim ~ s(Day, by = Group, k = 5)
1323 ##

```

```

1324 ## Parametric coefficients:
1325 ##           Estimate Std. Error t value Pr(>|t|)
1326 ## (Intercept)    21.929      1.518   14.45  <2e-16 ***
1327 ## ---
1328 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1329 ##
1330 ## Approximate significance of smooth terms:
1331 ##           edf Ref.df      F p-value
1332 ## s(Day):GroupControl  3.488  3.851  5.244 0.00442 **
1333 ## s(Day):GroupTreatment 2.962  3.461 24.272 < 2e-16 ***
1334 ## ---
1335 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1336 ##
1337 ## R-sq.(adj) =  0.502   Deviance explained = 53.5%
1338 ## -REML = 413.52   Scale est. = 230.47      n = 100

```

1340 Diagnostics for this model indicate that the k-index is still below 1 (0.32 from `gam.check`), and that the
1341 residuals are still not following a normal distribution (Figure 9). Moreover, the smooths (plotted via the
1342 `draw()` function) appear with a fairly linear profile, which indicates they are still not capturing the trends
1343 observed in the data. From `summary()`, the deviance explained by the model is $\approx 43\%$.

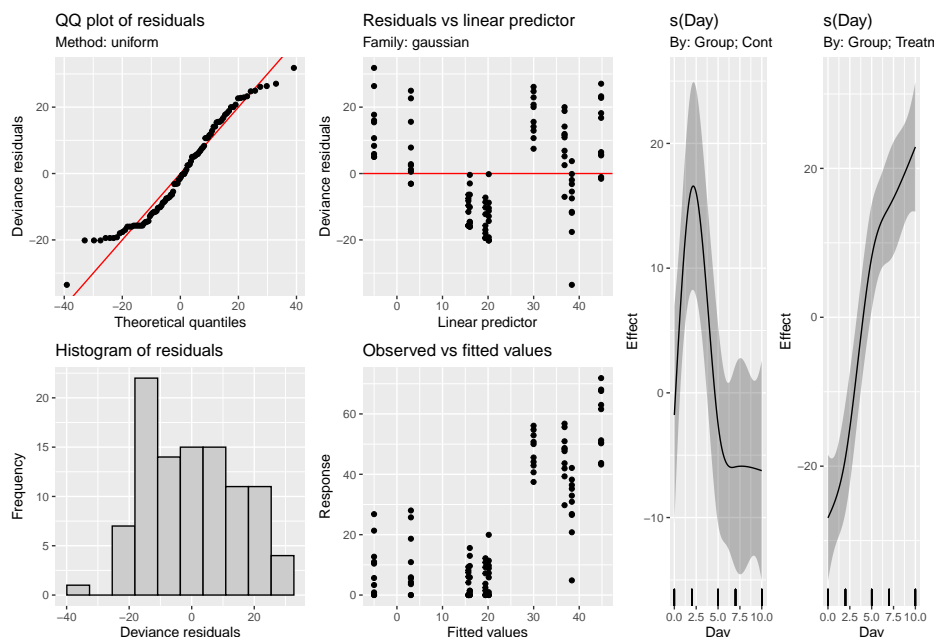


Figure 9: Graphical diagnostics for the second GAM model. Left: Graphical diagnostics provided by the function `appraise` from the package *gratia*. Right: Fitted smooth for the model, provided by the function `draw`.

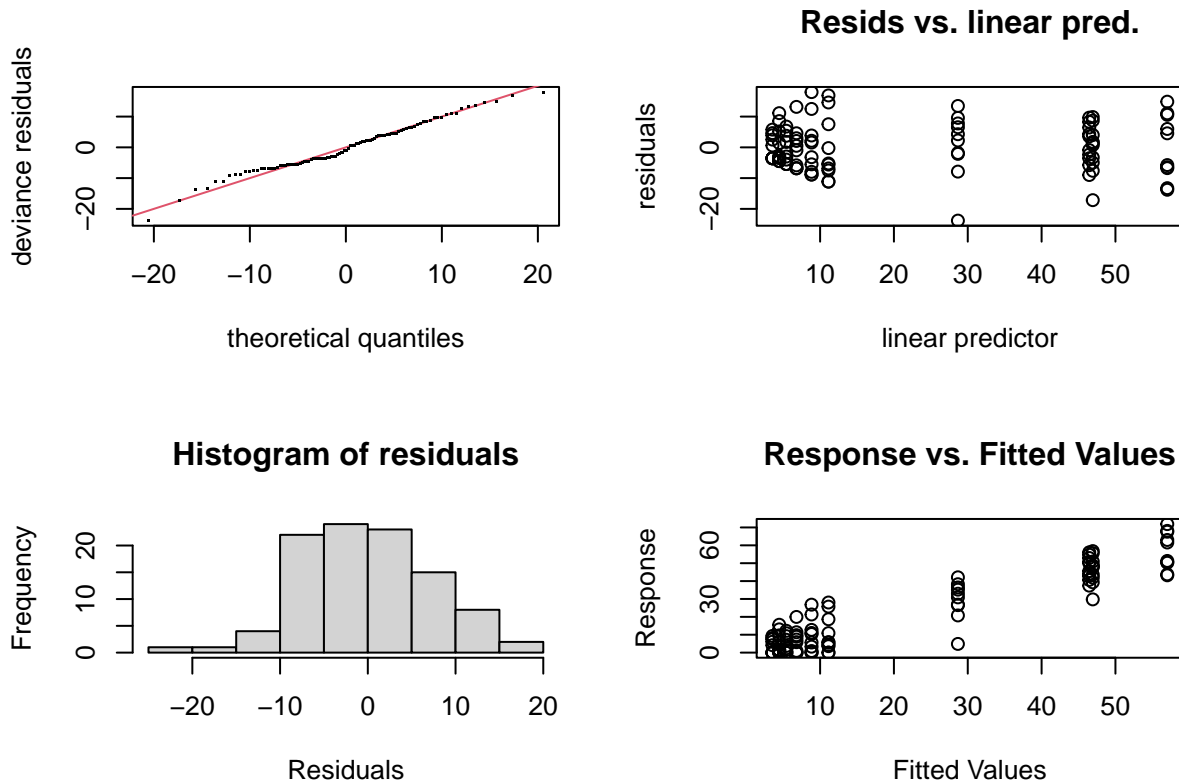
1344 B.1.3 Third model

1345 Model `gam_00` was built for didactic purposes to cover the simplest case, but it does not account for the
1346 nesting of the data by Group, which is apparent from the type of smooth fitted, the model diagnostics, and,
1347 the low variance explained by the model. On the other hand, `gam_01` takes into account the nesting within
1348 each group and provides better variance explanation, but as indicated in Section 5, in order to differentiate
1349 between each group a parametric term needs to be added to the model for the interaction of *Day* and *Group*.

```

1350 #GAM for StO2
1351
1352
1353 m1 <- gam(StO2_sim ~ Group+s(Day, by = Group, k = 5),
1354           method='REML',
1355           data = dat_sim)
1356
1357 gam.check(m1)
1358

```



```

1359
1360 ##
1361 ## Method: REML   Optimizer: outer newton
1362 ## full convergence after 9 iterations.
1363 ## Gradient range [-2.780424e-08,2.076237e-08]
1364 ## (score 354.6068 & scale 63.7304).
1365 ## Hessian positive definite, eigenvalue range [1.095531,48.08644].
1366 ## Model rank = 10 / 10
1367 ##
1368 ## Basis dimension (k) checking results. Low p-value (k-index<1) may
1369 ## indicate that k is too low, especially if edf is close to k'.
1370 ##
1371 ##
1372 ##           k'   edf k-index p-value
1373 ## s(Day):GroupControl  4.00 3.87   1.02   0.52
1374 ## s(Day):GroupTreatment 4.00 3.83   1.02   0.58
1375
1376
1377 summary(m1)
1378

```



```

1379 ##
1380 ##
1381 ## Family: gaussian
1382 ## Link function: identity
1383 ##
1384 ## Formula:
1385 ## St02_sim ~ Group + s(Day, by = Group, k = 5)
1386 ##
1387 ## Parametric coefficients:
1388 ##               Estimate Std. Error t value Pr(>|t|)
1389 ## (Intercept)      9.781      1.129   8.664 1.68e-13 ***
1390 ## GroupTreatment    24.296      1.597  15.217 < 2e-16 ***
1391 ## ---
1392 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1393 ##
1394 ## Approximate significance of smooth terms:
1395 ##               edf Ref.df    F p-value
1396 ## s(Day):GroupControl  3.867  3.989 19.38 <2e-16 ***
1397 ## s(Day):GroupTreatment 3.826  3.981 80.29 <2e-16 ***
1398 ## ---
1399 ## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
1400 ##
1401 ## R-sq.(adj) =  0.862   Deviance explained = 87.4%
1402 ## -REML = 354.61   Scale est. = 63.73       n = 100
1403

```

1404 The resulting model `ism1`, which is the model fitted in the main manuscript. By using `appraise()` and `draw`
1405 on this model (Figure 10) we see that the trend on the QQ plot has improved, the histogram of the residuals
1406 appears to be reasonably distributed, and the smooths are capturing the trend of the data within each group
1407 . From `gam.check`, the k-index is now at an acceptable value (≈ 1.02), and `summary` now indicates that the
1408 model is able to capture 87% of the variance data.

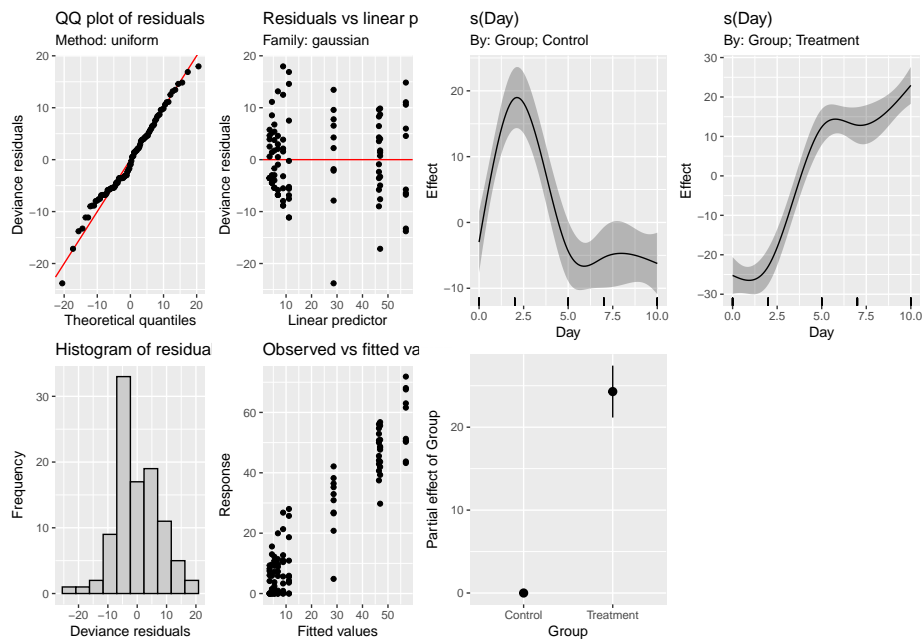


Figure 10: Graphical diagnostics for the final GAM model. Left: Graphical diagnostics provided by the function `appraise` from the package *gratia*. Right: Fitted smooths for the model, provided by the function `draw`.

B.1.4 Comparing models via AIC

One final comparison that can be made for model selection involves the use of the Aikake Information Criterion (AIC). This metric is used to estimate information loss, which we want to minimize with an appropriate model. Therefore, when 2 or more models are compared, the model with lower AIC is preferred. In R, the comparison is done using the `AIC` function.

```
AIC(gam_00, gam_01, m1)
```

	##		df	AIC
	##	gam_00	3.623938	891.2466
	##	gam_01	9.312053	838.6825
	##	m1	10.970436	710.9994

The output in this case is expected: model `gam1` has a lower AIC (711.46) whereas the initial two models have higher AICs (891 and 850). The AIC should not be considered as the only estimator of model goodness, instead to be used as complimentary information to the graphical diagnostics and model checks described above.

B.1.4.1 Pairwise comparisons of smooth confidence intervals The estimation of significant differences between each treatment group can be achieved via pairwise comparisons of the smooth confidence intervals as described in section 5.3. In this case, the “design matrix” is used to estimate the pairwise comparisons (see main manuscript for details and associated references). Briefly, the “design matrix” (also known as the “Xp matrix”) from the selected model (`m1`) is used to calculate a 95% confidence interval of the difference between the smooth terms for each group. This approach allows to estimate the time intervals where a significant difference exists between the groups (confidence interval above or below 0). **All pairwise comparisons in this paper have been centered at the response scale to ease interpretation .**

```
##Pairwise comparisons
pdat <- expand.grid(Day = seq(0, 10, length = 400),
                  Group = c('Control', 'Treatment'))

##matrix that contains the basis functions evaluated at the points in pdat
xp <- predict(m1, newdata = pdat, type = 'lpmatrix')

#Find columns in xp where the name contains "Control"
c1 <- grepl('Control', colnames(xp))

#Find columns in xp where the name contains 'Treatment'
c2 <- grepl('Treatment', colnames(xp))

#Find rows in pdat that correspond to either 'Control' or 'Treatment'
r1 <- with(pdat, Group == 'Control')
r2 <- with(pdat, Group == 'Treatment')

# In xp: find the rows that correspond to Control or Treatment, those that
do not match will be
#set to zero. Then, subtract the values from the rows corresponding
to 'Control' from those that correspond
#to 'Treatment'
X <- xp[r1, ] - xp[r2, ]

## remove columns that do not contain name 'Control' or 'Treatment'
```

```

1462 X[, ! (c1 | c2)] <- 0
1463 ## zero out the parametric cols, those that do not contain in the
1464 characters 's('
1465 #X[, !grepl('^s\\(', colnames(xp))] <- 0
1466
1467 #Multiply matrix by model coefficients. X has (p,n) (rows, columns)
1468 and the coefficient matrix has
1469 #dimensions (n,1). The resulting matrix has dimensions (p,1)
1470 dif <- X %%% coef(m1)
1471
1472 #comp<-test %%% coef(gam1)[3:10]
1473
1474 #Calculate standard error for the computed differences using the variance-
1475 covariance matrix
1476 #of the model
1477 se <- sqrt(rowSums((X %%% vcov(m1, unconditional = FALSE)) * X))
1478 crit <- qt(0.05/2, df.residual(m1), lower.tail = FALSE)
1479 #upper limits
1480 upr <- dif + (crit * se)
1481 #lower limits
1482 lwr <- dif - (crit * se)
1483 #put all components in a dataframe for plotting
1484 comp1<-data.frame(pair = paste('Control', 'Treatment', sep = '-'),
1485                   diff = dif,
1486                   se = se,
1487                   upper = upr,
1488                   lower = lwr)
1489
1490
1491
1492 #add time point sequence
1493 comp_St02 <- cbind(Day = seq(0, 10, length = 400),
1494                   rbind(comp1))
1495
1496 #plot the difference
1497 c1<-ggplot(comp_St02, aes(x = Day, y = diff, group = pair)) +
1498   #ribbon for difference confidence interval
1499   geom_ribbon(aes(ymin = lower, ymax = upper),
1500             alpha = 0.5,
1501             fill='#DB3A07FF') +
1502   geom_line(color='black',size=1) +
1503   geom_line(data=comp_St02,aes(y=0),size=0.5)+
1504   #highlight area under the curve where "Control" is higher
1505   geom_ribbon(data=comp_St02%%>%
1506             filter(lower>0),
1507             aes(ymin =0, ymax =lower),
1508             alpha = 0.5,
1509             fill='#30123BFF') +
1510   #highlight area under the curve where "Treatment" is higher
1511   geom_ribbon(data=comp_St02 %>%
1512             filter(upper<0),
1513             aes(ymin =0, ymax =upper),
1514             alpha = 0.5,
1515             fill='#7A0403FF') +

```

```

1516 facet_wrap(~ pair) +
1517 theme_classic()+
1518 labs(x = 'Days', y = expression(paste('Difference in StO2'[2] )))+
1519 scale_x_continuous(breaks=c(0,2,5,7,10))+
1520 theme(
1521   text=element_text(size=18),
1522   legend.title=element_blank()
1523 )
1524

```

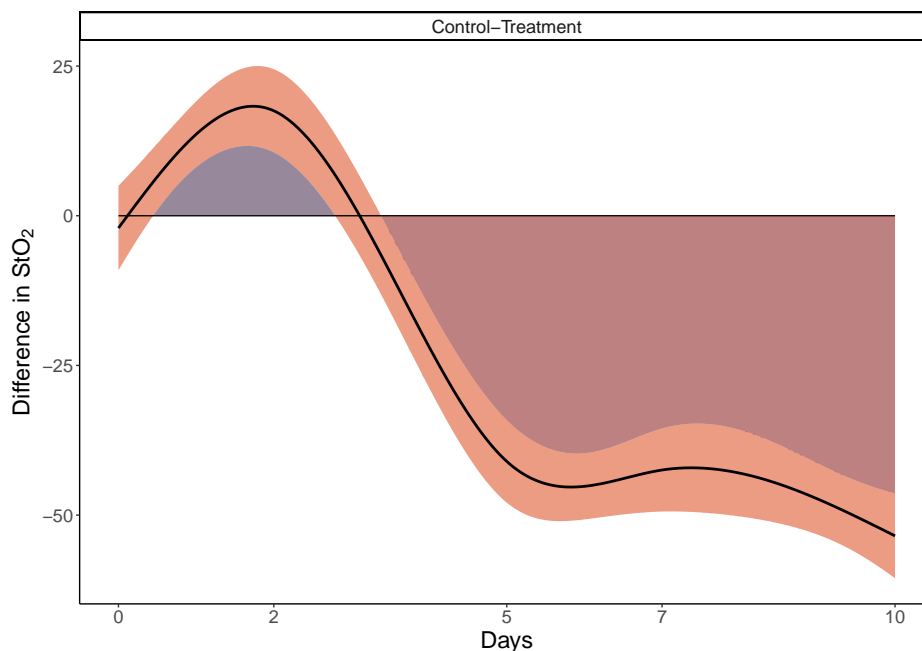


Figure 11: Smooth pairwise comparisons for model `m1` using a 95% confidence interval for the difference between smooths.

Of notice, a convenient wrapper for the function described above exists in the package `gratia`. In this package, `difference_smooths` is a function that makes the comparisons and produces Figure 11 when is used on a fitted model. The function syntax and an example can be found at:

<https://cran.r-project.org/web/packages/gratia/gratia.pdf>

Keep in mind that this function **does not** center the pairwise comparison at the response scale, so it has to be shifted in order to be compared to the raw data.

C GAM and Linear model plots and Missing data

This section covers the code used to generate Figure 3, where the simulated data, fit of the “final” GAM (`m1`), linear model and GAM on data with missing observations are presented. Note that panel A in Figure 3 and the inlet are generated in the code chunk where the data is simulated in Section B, and are called later to build the figure.

C.1 GAM and Linear model plots

This code chunk creates panels B and D in Figure 3. Note that this code uses the final GAM from the previous section (m1), so the simulated data and the model should be generated before running this section.

```
#linear model
lm1<-lm(St02_sim ~ Day + Group + Day * Group, data = dat_sim)

#creates a dataframe using the length of the covariates for the GAM
gam_predict <- expand_grid(Group = factor(c("Control", "Treatment")),
                          Day = seq(0, 10, by = 0.1),
                          subject=factor(rep(1:10)))

#creates a dataframe using the length of the covariates for rm-ANOVA
lm_predict<-expand_grid(Group = factor(c("Control", "Treatment")),
                        Day = c(0:10),
                        subject=factor(rep(1:10)),
                        )
lm_predict$subject<-factor(paste(lm_predict$subject, lm_predict$Group, sep
= "-"))

#adds the predictions to the grid and creates a confidence interval for
GAM
gam_predict<-gam_predict%>%
  mutate(fit = predict(m1,gam_predict,se.fit = TRUE,type='response')$fit
        ,
        se.fit = predict(m1, gam_predict,se.fit = TRUE,type='response')
        $se.fit)

#using lm
lm_predict<-lm_predict%>%
  mutate(fit = predict(lm1,lm_predict,se.fit = TRUE,type='response')$fit
        ,
        se.fit = predict(lm1, lm_predict,se.fit = TRUE,type='response')
        $se.fit)

#plot smooths and confidence interval for GAM
f3<-ggplot(data=dat_sim, aes(x=Day, y=St02_sim, group=Group)) +
  geom_point(aes(color=Group),size=1.5,alpha=0.5,show.legend = FALSE)+
  geom_ribbon(aes( x=Day,ymin=(fit - 2*se.fit),
                ymax=(fit + 2*se.fit),
                fill=Group
                ),
            alpha=0.3,
            data=gam_predict,
            show.legend=FALSE,
            inherit.aes=FALSE) +
  geom_line(aes(y=fit,
                color=Group),
            size=1,data=gam_predict,
            show.legend = FALSE)+
  #facet_wrap(~Group)+
  labs(y=expression(atop(St0[2], 'complete')))+
  scale_x_continuous(breaks=c(0,2,5,7,10))+
```

```

1590     theme_classic()+
1591     theme(
1592       axis.text=element_text(size=22)
1593     )+
1594     thm+
1595     thm1
1596
1597 #plot linear fit for rm-ANOVA
1598 f4<-ggplot(data=dat_sim, aes(x=Day, y=St02_sim, group=Group)) +
1599   geom_point(aes(color=Group),size=1.5,alpha=0.5,show.legend = FALSE)+
1600   geom_ribbon(aes( x=Day,ymin=(fit - 2*se.fit),
1601                 ymax=(fit + 2*se.fit),fill=Group),
1602             alpha=0.3,
1603             data=lm_predict,
1604             show.legend = FALSE,
1605             inherit.aes=FALSE) +
1606   geom_line(aes(y=fit,
1607                color=Group),
1608            size=1,data=lm_predict,
1609            show.legend = FALSE)+
1610   #facet_wrap(~Group)+
1611   labs(y=expression(paste('St0' [2], ' (simulated)')))+
1612   scale_x_continuous(breaks=c(0,2,5,7,10))+
1613   theme_classic()+
1614   theme(
1615     axis.text=element_text(size=22)
1616   )+
1617   thm+
1618   thm1
1619
1620
1621
1622 #posthoc comparisons for the linear model
1623 #library(multcomp)
1624
1625
1626 #summary(glht(lm1, linfct = mcp(Group = 'Tukey'))))
1627 #summary(glht(lm1, linfct=mcp(Group="Tukey", interaction_average=TRUE)))
1628

```

C.2 Working with Missing data in GAMs

This code chunk first randomly deletes 40% of the total observations in the original simulated data, and then an interaction GAM is fitted to the remaining data. Model diagnostics are presented, and an object that stores the fitted smooths is saved to be called in the final code chunk to build the figure.

```

1633 #missing data
1634 #create a sequence of 40 random numbers between 1 and 100, these numbers
1635   will
1636 #correspond to the row numbers to be randomly erased from the original
1637   dataset
1638
1639
1640 missing <- sample(1:100, 40)
1641

```

```

1642 #create a new dataframe from the simulated data with 40 rows randomly
1643     removed, keep the missing values as NA
1644
1645 ind <- which(dat_sim$StO2_sim %in% sample(dat_sim$StO2_sim, 40))
1646
1647 #create a new dataframe, remove the StO2 column
1648 dat_missing <- dat_sim[,-1]
1649
1650 #add NAs at the ind positions
1651 dat_missing$StO2_sim[ind]<-NA
1652
1653 #Count the number of remaining observations per day (original dataset had
1654     10 per group per day)
1655 dat_missing %>%
1656     group_by(Day,Group) %>%
1657     filter(!is.na(StO2_sim))%>%
1658     count(Day)

```

```

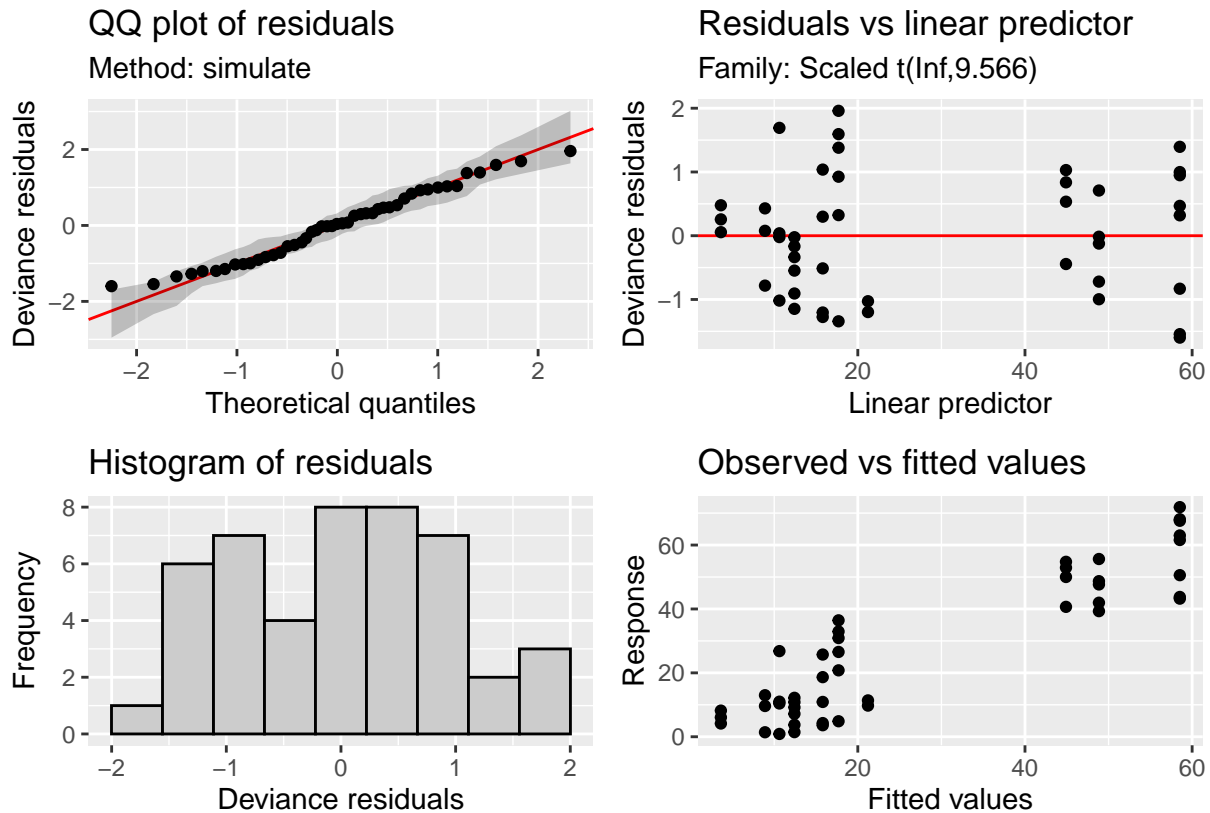
1660 ## # A tibble: 10 x 3
1661 ## # Groups:   Day, Group [10]
1662 ##   Day Group      n
1663 ##   <dbl> <fct>   <int>
1664 ## 1     0 Control     2
1665 ## 2     0 Treatment    4
1666 ## 3     2 Control     6
1667 ## 4     2 Treatment    5
1668 ## 5     5 Control     6
1669 ## 6     5 Treatment    4
1670 ## 7     7 Control     3
1671 ## 8     7 Treatment    5
1672 ## 9    10 Control     3
1673 ## 10    10 Treatment    8

```

```

1676 #the same model used for the full dataset
1677 mod_m1 <- gam(StO2_sim ~ Group+s(Day,by=Group,k=5), data = dat_missing,
1678     family=scat)
1679 #appraise the model
1680 appraise(mod_m1)

```



```

1683
1684
1685 m_predict <- expand_grid(Group = factor(c("Control", "Treatment")),
1686                           Day = seq(0, 10, by = 0.1))
1687
1688 #adds the predictions to the grid and creates a confidence interval
1689 m_predict<-m_predict%>%
1690   mutate(fit = predict(mod_m1,m_predict,se.fit = TRUE,type='response')$
1691     fit,
1692           se.fit = predict(mod_m1, m_predict,se.fit = TRUE,type='response
1693             ')$se.fit)
1694
1695
1696 f6<-ggplot(data=dat_missing, aes(x=Day, y=StO2_sim, group=Group)) +
1697   geom_point(aes(color=Group),size=1.5,alpha=0.5,show.legend = FALSE)+
1698   geom_ribbon(aes( x=Day,ymin=(fit - 2*se.fit),
1699                 ymax=(fit + 2*se.fit),
1700                 fill=Group
1701               ),
1702             alpha=0.3,
1703             data=m_predict,
1704             show.legend=FALSE,
1705             inherit.aes=FALSE) +
1706   geom_line(aes(y=fit,
1707                 color=Group),
1708             size=1,data=m_predict,
1709             show.legend = TRUE)+
1710   #facet_wrap(~Group)+

```



```

1711 labs(y=expression(atop(StO2[2], 'missing')))+
1712   scale_x_continuous(breaks=c(0,2,5,7,10))+
1713   theme_classic()+
1714   theme(
1715     axis.text=element_text(size=22)
1716   )+
1717   thm+
1718   thm1

```

```

1720
1721 mult_plot<-f2+inset_element(
1722   f1, left = 0.01,
1723   bottom = 0.5,
1724   right = 0.5,
1725   top = 1.0)+
1726   f3+f4+f6+
1727   plot_annotation(tag_levels='A')&
1728   ylim(c(-7,75)) &
1729   theme(
1730     text=element_text(size=18)
1731   )&
1732   thm
1733
1734 mult_plot
1735

```

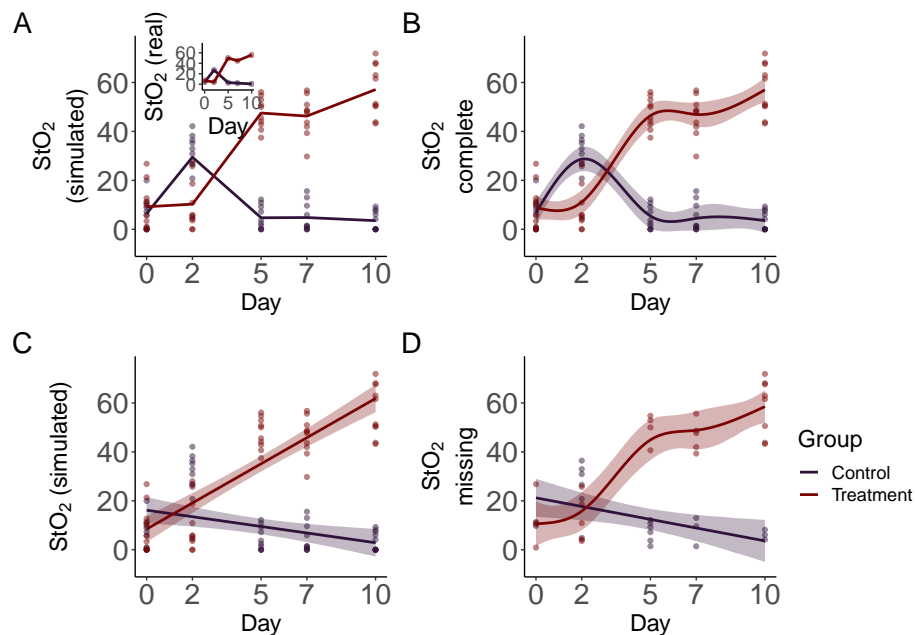


Figure 12: Simulated data and smooths for oxygen saturation in tumors. A: Simulated data that follows previously reported trends (inset) in tumors under chemotherapy (Treatment) or saline (Control) treatment. Simulated data is from a normal distribution with standard deviation of 10% with 10 observations per time point. Lines indicate mean oxygen saturation B: Smooths from the GAM model for the full simulated data with interaction of Group and Treatment. Lines represent trends for each group, shaded regions are 95% confidence intervals. C: rm-ANOVA model for the simulated data, the model does not capture the changes in each group over time. D: Smooths for the GAM model for the simulated data with 40% of its observations missing. Lines represent trends for each group, shaded regions are 95% confidence intervals.

C.3 Pairwise comparisons in GAMs: full and missing data cases

The next code chunk reproduces Figure 4. Here pairwise comparisons are made for the full and missing datasets.

```
##Pairwise comparisons

pdat <- expand.grid(Day = seq(0, 10, length = 400),
                  Group = c('Control', 'Treatment'))

#this function takes the model, grid and groups to be compared using the
  lpmatrix

smooth_diff <- function(model, newdata, g1, g2, alpha = 0.05,
                        unconditional = FALSE) {
  xp <- predict(model, newdata = newdata, type = 'lpmatrix')
  #Find columns in xp where the name contains "Control" and "Treatment"
  col1 <- grepl(g1, colnames(xp))
  col2 <- grepl(g2, colnames(xp))
  #Find rows in xp that correspond to each treatment
  row1 <- with(newdata, Group == g1)
  row2 <- with(newdata, Group == g2)
  ## difference rows of xp for data from comparison
  X <- xp[row1, ] - xp[row2, ]
  ## zero out cols of X related to splines for other lochs
  X[, ! (col1 | col2)] <- 0
  ## zero out the parametric cols
  #X[, !grepl('^s\\(', colnames(xp))] <- 0
  dif <- X %>% coef(model)
  se <- sqrt(rowSums((X %>% vcov(model, unconditional = unconditional))
    * X))
  crit <- qt(alpha/2, df.residual(model), lower.tail = FALSE)
  upr <- dif + (crit * se)
  lwr <- dif - (crit * se)
  data.frame(pair = paste(g1, g2, sep = '-'),
             diff = dif,
             se = se,
             upper = upr,
             lower = lwr)
}

comp1<-smooth_diff(m1,pdat,'Control','Treatment')

comp_StO2_full <- cbind(Day = seq(0, 10, length = 400),
                       rbind(comp1)) %>%
  mutate(interval=case_when(
    upper>0 & lower<0~"no-diff",
    upper<0~"less",
    lower>0~"greater"
  ))

c1<-ggplot(comp_StO2_full, aes(x = Day, y = diff, group = pair)) +
  geom_ribbon(aes(ymin = lower, ymax = upper),
            alpha = 0.5,
            fill='#DB3A07FF') +
```

```

1790 geom_line(color='#E75B64FF',size=1) +
1791 geom_line(data=comp_St02_full,aes(y=0),size=0.5)+
1792 geom_ribbon(data=comp_St02_full%>%
1793     filter(lower>0),
1794     aes(ymin =0, ymax =lower),
1795     alpha = 0.5,
1796     fill='#30123BFF') +
1797 geom_ribbon(data=comp_St02_full %>%
1798     filter(upper<0),
1799     aes(ymin =0, ymax =upper),
1800     alpha = 0.5,
1801     fill='#7A0403FF') +
1802 facet_wrap(~ pair) +
1803 theme_classic()+
1804 labs(x = 'Days', y = expression(paste('Difference in St0'[2] )))+
1805 scale_x_continuous(breaks=c(0,2,5,7,10))+
1806 theme(
1807     text=element_text(size=18),
1808     legend.title=element_blank()
1809 )
1810
1811
1812
1813 ###for missing data
1814 comp2<-smooth_diff(mod_m1,pdat,'Control','Treatment')
1815 comp_St02_missing <- cbind(Day = seq(0, 10, length = 400),
1816     rbind(comp2))
1817
1818 missing_plot<-ggplot(comp_St02_missing, aes(x = Day, y = diff, group =
1819     pair)) +
1820     geom_ribbon(aes(ymin = lower, ymax = upper), alpha = 0.2) +
1821     geom_line(color='black',size=1) +
1822     facet_wrap(~ pair) +
1823     labs(x = 'Days',
1824         y = expression(paste('Difference in St0'[2],'\n (missing data)'
1825             )))
1826     scale_x_continuous(breaks=c(0,2,5,7,10))+
1827     theme_classic()+
1828     theme(
1829         text=element_text(size=18),
1830         legend.title=element_blank()
1831     )
1832
1833 c2<-ggplot(comp_St02_missing, aes(x = Day, y = diff, group = pair)) +
1834     geom_ribbon(aes(ymin = lower, ymax = upper),
1835         alpha = 0.5,
1836         fill='#DB3A07FF') +
1837     geom_line(color='#E75B64FF',size=1) +
1838     geom_line(data=comp_St02_missing,aes(y=0),size=0.5)+
1839     # geom_ribbon(data=comp_St02_missing%>%
1840     #     filter(lower>0),
1841     #     aes(ymin =0, ymax =lower),
1842     #     alpha = 0.5,
1843     #     fill='#30123BFF') +

```

```

1844 geom_ribbon(data=comp_StO2_missing %>%
1845           filter(upper<0),
1846           aes(ymin =0, ymax =upper),
1847           alpha = 0.5,
1848           fill='#7A0403FF') +
1849 facet_wrap(~ pair) +
1850 theme_classic()+
1851 labs(x = 'Days', y = expression(paste('Difference in StO'2'[2] ))))+
1852 scale_x_continuous(breaks=c(0,2,5,7,10))+
1853 theme(
1854   text=element_text(size=18),
1855   legend.title=element_blank()
1856 )
1857
1858 pair_comp<-c1+c2

```

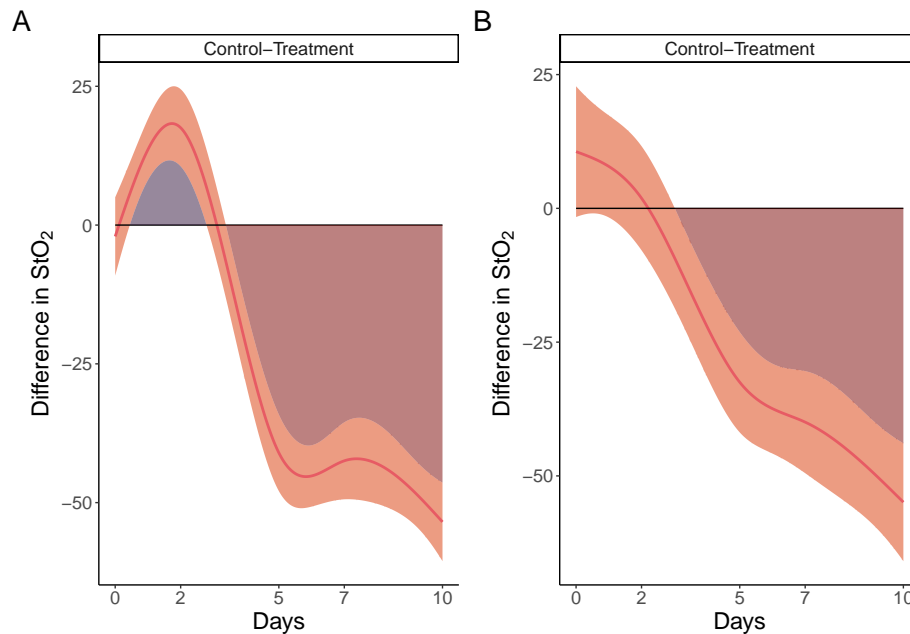


Figure 13: Pairwise comparisons for smooth terms. A: Pairwise comparisons for the full dataset. B: Pairwise comparisons for the dataset with missing observations. Significant differences exist where the interval does not cover 0. In both cases the effect of treatment is significant after day 5.